

pypresso

what it computes, and how to ask for it

Plane-wave DFT in Python and JAX, validated against Quantum ESPRESSO

August 31, 2026

Abstract

User documentation for the feature set: every capability, the equation behind it, a snippet that runs it, what it was validated against, and *what it refuses*. Two conventions run through the whole document.

Everything is validated against Quantum ESPRESSO on the same input, and the agreement is quoted per feature rather than claimed in general. Where a number is missing it is because the comparison does not exist, and that is said rather than hidden.

Anything not implemented is refused by name, with an error saying what is missing — never silently approximated. The “Refuses” notes are as much of the documentation as the features: they are the promise that a run which starts is a run whose physics is there.

One equation is written down and the rest are derived from it. Almost everything below is a derivative of the total energy, taken by the compiler rather than by hand:

$$E[\{\psi\}, \tau, \varepsilon] = T_s + E_H + E_{xc} + E_{\text{loc}} + E_{\text{nl}} + E_{\text{Ewald}}$$

quantity	what it is	how it is obtained
$v_{xc}(\mathbf{r})$	$\delta E_{xc}/\delta n$	jax.grad of the energy density
$\mathbf{F}_I$	$-\partial E/\partial \boldsymbol{\tau}_I$	jax.grad at frozen ψ
σ_{ij}	$-\Omega^{-1}\partial E/\partial \varepsilon_{ij}$	jax.grad along a strain
$C_{I\alpha, J\beta}$	$\partial^2 E/\partial \tau_{I\alpha} \partial \tau_{J\beta}$	jvp of the force
$Z_{I, \alpha\beta}^*$	$\partial F_{I\alpha}/\partial \mathcal{E}_\beta$	jvp of the force along the field response
$\partial \epsilon_{ij}/\partial \tau$	Raman	jvp of the second-order energy

The “frozen ψ ” is what makes this exact rather than approximate: at self-consistency the energy is stationary with respect to the wavefunctions, so the terms from $\partial\psi/\partial\lambda$ vanish. That is the Hellmann-Feynman argument, and taking it one order up is the $2n+1$ theorem, which is why a *third* derivative needs only first-order wavefunctions.

Contents

1	Input and invocation	3
2	Ground state	5
3	Bands, densities of states, projections	5
4	Pseudopotentials and functionals	6
5	Spin, magnetism and spin-orbit	7
6	Forces, stress and relaxation	8
7	Response: dielectric, phonons, Raman	9
7.1	The piezoelectric tensor	12
7.2	Spin-polarized response	13
8	Optical spectra and excitons: TDDFT	13
8.1	The conductivity tensor, the Kerr angle and the anomalous Hall effect . . .	15
9	Topological invariants	16
9.1	The curvature as a map: the Kubo route	17
10	Effective masses and site angular momenta	18
10.1	The effective mass tensor	18
10.2	$\langle L \rangle$, $\langle S \rangle$ and $\langle J \rangle$ on each atom	19
11	Things with no pw.x counterpart	19
12	Performance: dials, memory and GPUs	20
13	Accuracy summary	21
	Where to go next	21

1 Input and invocation

pypresso reads **pw.x input files**. The same file that runs in Quantum ESPRESSO runs here, which is what makes every comparison in this document possible.

```
from pypresso import Calculator

calc = Calculator.from_file("scf.in", pseudo_dir="pseudo", conv_thr=1.0e-8)
result = calc.get_scf()

print(result.total_energy, "Ry in", result.iterations, "iterations")
```

`conv_thr` means what it means in a `pw.x` input: it is compared against QE's `dr2`, the Hartree energy of the density residual.

It need not be given at all. `from_file` and `from_text` **adopt the input's own &electrons namelist** — `conv_thr`, `mixing_beta`, `mixing_mode`, `mixing_fixed_ns` and `electron_maxstep` — so an input file that states how to converge itself converges the same way here as under `pw.x`. A keyword argument given to the constructor still wins over the file, which is how the snippet above overrides it. One variable of that namelist is read by `pw.x` and deliberately *not* adopted: `diagonalization = 'cg'` would turn a run that works into an error, and mapping it silently onto Davidson would be the kind of substitution refused everywhere else here. Name a solver at construction instead.

One object, bound methods

A `Calculator` is a system together with its pseudopotentials, and every calculation in this document is a method on it. The pseudopotentials are read from the names the `ATOMIC_SPECIES` card gave, resolved against `pseudo_dir` — which `from_file` defaults to the input file's own directory — so a script names an input and then names physics:

```
calc = Calculator.from_file("scf.in", pseudo_dir="pseudo")

bands = calc.get_bands(kpoints=path)      # runs the SCF first if none is cached
dos = calc.get_dos(grid=(8, 8, 8))
forces = calc.get_forces()
epsilon = calc.get_dielectric_tensor()
phonons = calc.get_phonons()

bands.plot()                             # and each result draws itself
```

Three properties are worth knowing before relying on it.

The ground state is cached, and computed on demand. A method that needs one and finds no cache runs an SCF, printing a line to `stderr` saying so; `get_scf()` is the explicit path and is the same code. The cache is a *single slot* holding the result together with the options that produced it — calling `get_scf` again with different options reruns and replaces it — because what it holds is the wavefunctions, and a cache keyed by option sets would quietly hold several sets of them. `calc.scf_result` reads the slot without triggering anything, and a state that did not converge is refused rather than differentiated.

Nothing mutates. `with_positions`, `with_cell` and `with_spin` return a *new* calculator with an

empty cache; the converged state crosses as a starting guess where that is sound, which for `with_spin` is the spin-regime continuation described under *Spin, magnetism and spin-orbit* and converges in one iteration rather than twenty-five. A cached result under a moved atom would otherwise answer for the geometry it was converged at.

The refusals pass through untouched. Every restriction in this document — PAW Born charges, a metal’s Raman tensor, a spiral with spin-orbit coupling — raises through the bound method exactly as through the function beneath it.

The functional entry points the rest of this document names — `run_scf`, `run_bands`, `dielectric_tensor` and the others — are unchanged and remain the way to drive a calculation whose state is being managed by hand:

```
from pypresso.io.pwin import read_pw_input
from pypresso.pseudo import read_upf
from pypresso.scf.driver import Calculation, run_scf
from pypresso.system import build_system

system = build_system(read_pw_input("scf.in"))
pseudos = tuple(read_upf(f"pseudo/{s.pseudo_file}") for s in system.structure.species)
result = run_scf(system, pseudos, conv_thr=1.0e-8)
```

What such a call has to carry with the density is the rest of the converged state — `becsum` for a PAW dataset, `ns` under a Hubbard U , `tau` under a meta-GGA. None of the three can be rebuilt from the density, and the entry points refuse without them rather than computing something else. The calculator passes them.

A command line covers the common workflows:

```
python3 -m pypresso.cli inspect <qe-output> # summarise what the parser reads
python3 -m pypresso.cli dos <input> # density of states
python3 -m pypresso.cli pdos <input> # projected DOS
python3 -m pypresso.cli relax <input> # structural relaxation
python3 -m pypresso.cli stress <input> # stress tensor
python3 -m pypresso.cli spiral <input> # spin-spiral scan
```

Standard `pw.x` variables — `ecutwfc`, `ecutrho`, `ibrav`, `nbnd`, `occupations`, `degauss` and the rest — mean what they mean there and are not re-documented here. A handful of knobs are worth knowing because they are specific to this code or to a feature below:

<code>mbj_c</code>	fixes Tran-Blaha’s c instead of taking the cell average
<code>london_s6</code> , <code>london_rcut</code>	Grimme D2’s scaling and real-space cutoff
<code>cell_dofree</code>	which cell degrees of freedom a vc-relax may move
<code>starting_ns_eigenvalue</code>	seeds the DFT+U occupation matrix
<code>spiral_q</code>	the spin-spiral wavevector, in lattice coordinates
<code>LOCAL_MAGNETIC_FIELDS</code>	a card with no <code>pw.x</code> counterpart

Units are Rydberg atomic units throughout (Ry, bohr), matching QE. The only layer that speaks anything else is `io/`: Hubbard U in the HUBBARD card is in eV and is converted at the boundary, as `pw.x` does it.

Refuses. `K_POINTS gamma` selects the half-sphere storage of the gamma-point trick, which is generated but not consumed. Such a run is substituted by an explicit $k = 0$ on the full sphere — the same physics at twice the storage — and it says so.

2 Ground state

The Kohn-Sham problem, generalised because ultrasoft and PAW make the overlap non-trivial:

$$\begin{aligned}\hat{H}|\psi_{n\mathbf{k}}\rangle &= \epsilon_{n\mathbf{k}} \hat{S}|\psi_{n\mathbf{k}}\rangle, & \hat{S} &= 1 + \sum_{I,ij} q_{ij} |\beta_i^I\rangle\langle\beta_j^I| \\ \hat{H} &= -\nabla^2 + v_{\text{loc}}(\mathbf{r}) + v_{\text{H}}[n] + v_{xc}[n] & & + \sum_{I,ij} D_{ij}^I |\beta_i^I\rangle\langle\beta_j^I|\end{aligned}$$

with $n(\mathbf{r}) = \sum_{n\mathbf{k}} w_{\mathbf{k}} f_{n\mathbf{k}} (|\psi_{n\mathbf{k}}|^2 + \sum_{ij} Q_{ij}^I \langle\psi|\beta_i^I\rangle\langle\beta_j^I|\psi\rangle)$, the augmentation term being what ultrasoft adds. D_{ij}^I is rebuilt from the potential every iteration, which is why the SCF is a fixed point in (n, D) and not in n alone. Convergence is measured as QE measures it — the Hartree energy of the residual:

$$dr^2 = \iint \frac{\delta n(\mathbf{r}) \delta n(\mathbf{r}')}{|\mathbf{r} - \mathbf{r}'|} d\mathbf{r} d\mathbf{r}' < \text{conv_thr}$$

```
result = run_scf(system, pseudos, conv_thr=1.0e-10, mixing_mode="anderson")
print(result.total_energy, result.energy_terms) # QE's own term-by-term split
print(result.fermi_energy, result.homo, result.lumo)
```

Metals work with every smearing QE has (gaussian, marzari-vanderbilt, methfessel-paxton, fermi-dirac) and with the tetrahedron family — both linear and Bloechl-corrected — usable as an occupation scheme *inside* the SCF, not only for a density of states.

Continuation across a change of spin regime. Give `run_scf` a previous result as `starting_from` — `System.with_spin` underneath — and it promotes a converged state into a target's variables: non-magnetic into collinear, collinear into noncollinear, spin-orbit switched on. It reaches the *same* solution as a fresh run ($\leq 4\text{e-}8$ Ry on six cases) in as little as one iteration. The magnetization is **seeded** from the target's `starting_magnetization` when the source has none, because nothing in the SCF breaks spin symmetry on its own and an unseeded promotion would converge back to the unpolarized solution and report success.

Convergence. `mixing_mode` selects `anderson` (default), `linear`, `broyden`, or Kerker/Thomas-Fermi preconditioning (TF); `scf_solver` offers a Newton-Krylov alternative that is slower but reaches solutions the mixer cannot hold. `max_iterations` defaults to 100 — **a hard SCF may need more**, and hitting the cap is reported as not converged rather than as an answer.

3 Bands, densities of states, projections

A band structure is one diagonalisation per $\mathbf{k}$ on a fixed density — `run_bands`, or `run_nscf` for a non-self-consistent run on a grid rather than a path. A density of states is

$$g(E) = \sum_{n\mathbf{k}} w_{\mathbf{k}} \delta(E - \epsilon_{n\mathbf{k}})$$

with δ replaced by a smearing function or evaluated by tetrahedron interpolation. The projected DOS weights each term by $|\langle\phi_\mu|\hat{S}|\psi_{n\mathbf{k}}\rangle|^2$, and the spilling parameter is what is *not* captured, $\sum_{n\mathbf{k}} w_{\mathbf{k}} f_{n\mathbf{k}} (1 - \sum_\mu |\langle\phi_\mu|\hat{S}|\psi_{n\mathbf{k}}\rangle|^2)$.

```

from pypresso.workflows.bands import run_bands
from pypresso.workflows.dos import run_dos
from pypresso.workflows.pdos import run_pdos

bands = run_bands(system, pseudos, scf.density, kpoints=path)  # a k-path

dos, _ = run_dos(system, pseudos, scf.density, grid=(12, 12, 12),
                 scheme="tetrahedra_opt")  # -> (DensityOfStates, NSCFResult)

pdos, _ = run_pdos(system, pseudos, scf)  # -> (ProjectedDOS, NSCFResult)
print(pdos.channels[0])  # resolved by atom, l and m
print(pdos.charges.charges, pdos.charges.spilling)

```

The DOS schemes live behind a name registry (DOS_SCHEMES), and the projected DOS feeds *the same* registry as a per-band weight, so a PDOS and a DOS are the same integration. Projections are $\langle \phi | \hat{S} | \psi \rangle$ on Löwdin-orthogonalised pseudo-atomic orbitals. Validated against a purpose-built projwfc.x on seven cases: **6.9e-4** on a projection, **4.7e-5** on a Löwdin charge, and the spilling parameter to all four decimals it prints.

4 Pseudopotentials and functionals

Kinds: norm-conserving, **ultrasoft** and **PAW** — the two-grid split, the augmentation charge, the overlap operator, self-consistent D_{ij} , and PAW’s one-centre terms including its radial Poisson solve and spherical quadrature. UPF v2 is read.

Functionals: LDA (PZ), **PBE**, **revPBE**, **PBEsol**, on the plane-wave grid and on the PAW spheres. The functional comes from the pseudopotential headers unless input_dft overrides it. Only the *energy* is written down; v_{xc} , and a GGA’s v_1 and v_2 , come from jax.grad of it.

Meta-GGA, potential-only branch: input_dft = 'tb09' (Tran-Blaha) and 'bj06' (Becke-Johnson). These invert the rule above — they *are* potentials, there is no energy — so the consequences are enforced rather than documented: run_scf warns that its total is not the value of any functional it minimised, and forces, stress, phonons and response all refuse. Silicon’s gap goes from LDA’s 0.49 eV to **1.13 eV** against an experimental 1.17.

Van der Waals: Grimme **D2** (vdw_corr = 'grimme-d2'). Bilayer graphene matches pw.x to **3.1e-9 Ry** and binds at 3.23 Å where PBE alone has no minimum.

Refuses. UPF v1; an unimplemented functional (rather than silently substituting LDA); energy-carrying meta-GGAs (TPSS, SCAN, M06L), whose potential has a $\partial E / \partial \tau$ piece acting on the wavefunction that is not written; plain *ultrasoft* under a meta-GGA, which has no partial waves to reconstruct τ from inside the sphere where PAW does; EXX; and **D3, Tkatchenko-Scheffler, MBD and XDM** — where pw.x warns and silently runs with no correction at all.

5 Spin, magnetism and spin-orbit

regime	how to ask	note
unpolarised	default	
collinear	nspin = 2	one Fermi level, or two under tot_magnetization
noncollinear	noncolin = .true.	magnetization is a vector
spin-orbit	+ lspinorb = .true.	<i>j</i> -resolved projectors; NC, US and PAW

Three spin numbers are kept apart, and `System` exposes all three: `nspin` says which regime is in force, `npol` is the number of spinor components of a *wavefunction*, and `nspin_mag` the number of components of a *density*. They coincide at 1 and 2 and come apart at 4, where `npol` = 2 but `nspin_mag` is **one** for a nonmagnetic spin-orbit run — which is why such a run costs about what a doubled unpolarized one costs.

Fixed occupations under `nspin` = 2 fill each channel to its own count rather than searching for a level: `tot_magnetization` gives `nelup` and `nelow`, and `iweights_only` fills `NINT( $n_{\uparrow}$ )` bands in one channel and `NINT( $n_{\downarrow}$ )` in the other. That is the *only* shape the combination has — `pw.x` refuses fixed-occupation LSDA without a `tot_magnetization`, and requires an integer one — and this package makes the same two refusals at input, before anything is allocated. Each channel's highest occupied level is reported as `fermi_energy_up` / `fermi_energy_down` beside the overall `homo` and `lumo`, which is what `iweights` returns. Validated on an oxygen atom at `tot_magnetization` = 2, whose four-up / two-down filling is the Hund's-rule ground state: **4.9e-9 Ry** against `pw.x`.

Magnetic fields and constraints: a uniform field over the cell, a field inside one atom's sphere (through a `LOCAL_MAGNETIC_FIELDS` card that `pw.x` has no counterpart for), Elk's `reducebf`, and all four of QE's `constrained_magnetization` schemes. **The field's energy is not part of the reported total** — QE prints `etcon` and never adds it, and Elk excludes its external field's energy by the same convention.

DFT+U: QE's `lda_plus_u_kind` = 0 (Dudarev) with the J_0/β extension, on atomic, ortho-atomic and norm-atomic projectors, read from the `HUBBARD` card. Forces come from differentiating through projectors that move with the atoms.

Refuses. The full (Liechtenstein) DFT+U formulation, intersite V , background channels, the orbital-resolved variant, the `wf/pseudo` projector sets, and noncollinear `ns_nc`; and PAW + GGA + a noncollinear magnetization together.

A **magnetic field in a noncollinear run with no starting moment** is refused rather than run. `domag` comes from `starting_magnetization` and from nothing else (`setup.f90`), so such a run has `nspin_mag` = 1 and a density with no magnetization channels for the field to act on. Set `starting_magnetization` for one species — any nonzero value will do, the field decides where the moment goes. Worth knowing: `pw.x` *accepts* this input and silently drops the field, because `add_bfield` writes it into the potential's magnetization channels and `vloc_psi_nc` applies those only IF (`domag`).

The **residual solver** cannot take a fixed occupation whose filled/empty boundary cuts a **degenerate multiplet** — which is exactly what an atom's Hund's-rule filling does. Which member of the multiplet the eigensolver returns is arbitrary, so the density map is not a function of the density and a Newton step has nothing to converge on; it is diagnosed by name rather than left as a NaN. The mixer damps its way past it, and a *gapped* fixed-occupation LSDA cell agrees between the two solvers to 5e-12 Ry.

6 Forces, stress and relaxation

$$\mathbf{F}_I = -\frac{\partial E}{\partial \tau_I} \Big|_{\psi \text{ fixed}}, \quad \sigma_{ij} = -\frac{1}{\Omega} \frac{\partial E}{\partial \varepsilon_{ij}} \Big|_{\psi \text{ fixed}}$$

Both are one gradient. With ultrasoft the orthonormality constraint $\langle \psi_n | \hat{S} | \psi_m \rangle = \delta_{nm}$ is carried explicitly, so its multiplier term — QE’s Pulay contribution — falls out of the same derivative rather than being added afterwards. A variable-cell relaxation minimises the **enthalpy**, which is why a relaxed crystal carries the applied pressure rather than having zero stress:

$$H = E + P \Omega, \quad \frac{\partial H}{\partial h} = \Omega (P \mathbf{1} - \sigma) h^{-T}$$

```
from pypresso.forces import compute_forces
from pypresso.stress import compute_stress
from pypresso.workflows.relax import run_relax
from pypresso.workflows.vc_relax import run_vc_relax

forces = compute_forces(calculation, result) # method=... for QE's own terms
print(forces.forces) # (nat, 3) in Ry/bohr

stress = compute_stress(calculation, result, terms=True)
print(stress.pressure) # kbar

relaxed = run_relax(system, pseudos, forc_conv_thr=1.0e-4)
cell = run_vc_relax(system, pseudos, press=500.0) # kbar
print(cell.pulay_error) # the frozen-basis error, reported not hidden
```

QE’s `force_lc`, `force_cc`, `force_ew`, `force_us`, `addusforce` and `force_corr` are transcribed too, behind the same registry — **as a cross-check, not as the implementation**. The two share no machinery, which is what makes disagreement informative; it is how the augmentation force’s sign and a missing gradient correction were found.

For a noncollinear or spin-orbit run the same expression is evaluated on two-component spinors: the nonlocal term takes the pseudopotential’s bare `dvan_so` (a complex 2×2 matrix in spin space) and the orthonormality constraint takes `qq_so`, and the state is one coefficient vector of length $2n_{\text{pw}}$ rather than two. Nothing else in the functional changes.

```
from pypresso import Calculator

calculator = Calculator.from_file("pt2-soc-force.in") # noncolin, lspinorb
forces = calculator.get_forces() # (nat, 3) Ry/bohr
stress = calculator.get_stress() # (3, 3) Ry/bohr^3
print(forces.max_force, stress.pressure_kbar)
```

Validated against `pw.x` with `tpnfor/tstress` on a four-atom noncollinear hydrogen chain (norm-conserving) to 8.9×10^{-7} Ry/bohr, doubled fcc platinum with spin-orbit coupling to 7.5×10^{-6} (ultrasoft) and 7.3×10^{-7} (PAW), and the stress on those three plus QE’s own three `pw_spinorbit` cases to $\leq 1.2 \times 10^{-6}$ Ry/bohr³; and, without any Fortran at all, against a central finite difference of the frozen energy to 6.2×10^{-9} Ry/bohr.

A noncollinear or spin-orbit run has forces, a stress and a relaxation, but only through the *autodiff* method: `method='analytic'` is refused by name, because `force_us/stres_knl`

are transcriptions of QE's hand-derived expressions and have no spinor form. The force on an atom of a **spin spiral** is refused separately — its two spinor components live on different plane-wave spheres, so the nonlocal term needs the projectors of both; dE/dq is what a spiral has instead. **Elastic constants and electrostriction** stay refused for `noncolin`: they differentiate the same functional a third time and their first-order wavefunctions come from a Sternheimer solve that has no spinor form.

7 Response: dielectric, phonons, Raman

The **velocity operator** comes first and everything else uses it: it is one jvp of $\hat{H}(\mathbf{k})$ at a frozen sphere, because $\partial\hat{H}/\partial k_\alpha = i[\hat{H}, r_\alpha]$ in the periodic gauge. `band_velocities` exposes it — `Calculator.get_band_velocities` is the short way — and because the overlap carries a velocity too, a band velocity is $\langle\psi|\partial_k\hat{H} - \epsilon\partial_k\hat{S}|\psi\rangle$.

```
calc = Calculator.from_file("scf.in", pseudo_dir="pseudo")

velocities = calc.get_band_velocities()          # .velocities is (nk, nbnd, 3)
elsewhere = calc.get_band_velocities(kpoints=path)
```

Every first-order quantity then comes from the Sternheimer equation — the response of an occupied orbital to a perturbation, projected outside the occupied manifold so that no empty states and no energy denominators appear:

$$(\hat{H} - \epsilon_n\hat{S} + \alpha\hat{Q})|\Delta\psi_n\rangle = -\hat{P}_c^+ \Delta V|\psi_n\rangle$$

It is solved self-consistently, because ΔV contains the response of the potential to the response of the density,

$$\Delta V_{\text{scf}} = \Delta V_{\text{ext}} + \int K(\mathbf{r}, \mathbf{r}') \Delta n(\mathbf{r}') d\mathbf{r}'$$

and the kernel K is one jvp of `v_of_rho` rather than a second implementation. From its solutions:

$$\epsilon_{\alpha\beta}^\infty = \delta_{\alpha\beta} + \frac{8\pi}{\Omega} \sum_{\mathbf{k}} w_{\mathbf{k}} \langle \Delta_\alpha \psi_n | \partial_\beta \psi_n \rangle, \quad Z_{I,\alpha\beta}^* = \frac{\partial F_{I\alpha}}{\partial \mathcal{E}_\beta}$$

$$C_{I\alpha,J\beta} = \frac{\partial^2 E}{\partial \tau_{I\alpha} \partial \tau_{J\beta}}, \quad \omega^2 \mathbf{e} = \frac{1}{\sqrt{M_I M_J}} C \mathbf{e}, \quad \frac{\partial \epsilon_{ij}}{\partial \tau_{I\gamma}} \text{ (Raman)}$$

The Raman tensor is the second-order energy differentiated once more along a displacement, at *frozen first-order* wavefunctions — the $2n+1$ theorem, which is why a third derivative costs one more jvp and not a new solve.

```
from pypresso.response import (dielectric_tensor, dynamical_matrix,
                               raman_tensors, vibrational_spectrum)

eps = dielectric_tensor(calculation, scf.wavefunctions, scf.eigenvalues,
                        scf.density)
print(eps.epsilon)          # 3x3, cartesian

ph = dynamical_matrix(calculation, scf.wavefunctions, scf.eigenvalues,
```

```

scf.density, scf.becsum)
print(ph.frequencies)          # cm-1

raman = raman_tensors(calculation, scf, born_charges=True, keep_internals=True)
spectrum = vibrational_spectrum(calculation, scf) # per-mode Raman and IR
print(spectrum.raman_activity, spectrum.depolarisation)

```

born_effective_charges gives Z^* on its own, and mode_activities is the bare assembly — a transcription of dynmat.x's RamanIR, kept a pure function of arrays so it is testable without a self-consistent run.

Ultrasoft and PAW work for the dielectric constant ($\leq 1.2\text{e-}4$ against ph.x on four cases) *and for the dynamical matrix* — silicon's optical mode comes out at **513.295** and **513.378** cm^{-1} against ph.x's 513.275 and 513.404, which is tighter than the norm-conserving case's 0.05. Four things make the difference and all four switch themselves off when S is the identity: the source term is $(dH/du - \epsilon dS/du)|\psi\rangle$; the first-order state has an occupied block the Sternheimer solve does not produce; the mixed state changes at *frozen* states, because the augmentation charge and the projectors travel with their atom; and the orthonormality multipliers are variables of the functional that move — as a **matrix**, since a diagonal one is not invariant under the occupied-manifold rotation the state tangent is free in. Metals work for the response and for a norm-conserving dynamical matrix; and a symmetry-reduced wedge works as of the rank- n symmetriser.

The Raman tensor is ultrasoft and PAW as well, at **1.2e-4** on both against a central difference of ϵ over re-converged displaced cells, where the norm-conserving control on the same script is $6.8\text{e-}4$. It took **two tangents that are only right together**, and either alone is worse than neither: the state tangent is $P_c d\psi + \psi^{\text{ort}}$, the occupied block again; and $|b\rangle = P_c \mathbf{r}|\psi\rangle$ is *not* the solution of its own linear equation once adddvpsi_us has applied S to it and added the augmentation dipole, so db is the tangent of a *composition* rather than of a solve. What located them is that $\partial\epsilon/\partial\tau$ is a sum of five partial derivatives and each one can be finite-differenced on its own against a re-converged run: three agreed to $7\text{e-}4$ and two did not.

A trap worth knowing: a response on a reduced k -set is a *polar vector field* and must be symmetrised as one. Running the whole grid instead is sound only if the grid is closed under the point group — a **shifted** Monkhorst-Pack grid is not.

Refuses. Born charges for PAW (int3_paw against becsumort is missing, and without it the expression is wrong in sign as well as size); $\chi^{(2)}$ and the electro-optic tensor (the second-order source term has nothing to build the $\langle u_i | r_k | u_j \rangle$ piece from — and it is **42%** of the answer, measured); the non-analytic LO-TO term; phonons at $q \neq 0$; the dynamical matrix of an ultrasoft or PAW *metal* (the wg/wk weight split was derived for a response whose becsum dependence is entirely inside dps_i, and an insulator cannot say which weight the three further tangents belong with); any response under a potential-only meta-GGA; and a shifted grid with a reduced k -set.

Elastic constants, electrostriction and the elasto-optic tensor

The same machinery in the *strain* coordinate rather than the atomic one. **The strain response itself works on all three pseudopotential kinds** (strain_response): the augmentation charge $Q_{ij}(\mathbf{r})$ is a function of the *cell*, so a homogeneous strain deforms the table it is tabulated on where a displacement only translates it, and on top of that come the four terms the dynamical matrix needed. Against a central difference of *re-converged*

strained runs — which needs no `ph.x`, since `ph.x` has no strain perturbation — it agrees to **4.6e-4** (ultrasoft) and **4.7e-4** (PAW) against a norm-conserving 1.9e-4 on the same cell. The *derivatives* below are norm-conserving still — but the variational second-order energy they all stand on is not: it reproduces `dielec.f90`’s dielectric constant to **3.4e-10** (ultrasoft) and **6.9e-11** (PAW) against a norm-conserving 8.4e-10, which took four terms that each vanish when S is the identity. One of them is worth stating on its own, because it is a distinction rather than a term: a *state* is projected with $1 - \sum |\psi\rangle\langle\psi|S$ and a *right-hand side* with $1 - \sum S|\psi\rangle\langle\psi|$, and collapsing the two — which is exact at $S = 1$ — costs a factor of fourteen. The third derivative in this *strain* coordinate is what stays norm-conserving, and how far off it is has been measured rather than left unknown: both of the tangents that closed the Raman tensor transfer, and take it from **4.6e-2** to **1.3e-2** against a central difference over re-converged strained cells, where the norm-conserving control on the same script is **2.3e-4**. What is left is entirely the *db* term — the same -1.72 of 112 on ultrasoft and on PAW, which is what says it is structural. One candidate for it is excluded by measurement: writing the commutator that sources *db* with the multiplier matrix rather than the frozen scalar eigenvalue closes this coordinate (1.7e-4) and breaks the *displacement* one, taking the Raman tensor from 1.2e-4 to 1.14e-3. `strain_response` solves the first-order problem for a strain — Abinit’s metric-tensor formulation, obtained for nothing here because the cell was already written in reduced coordinates — and two things follow from it:

$$C_{ijkl} = \frac{\partial \sigma_{ij}}{\partial \varepsilon_{kl}}, \quad \frac{\partial \chi_{ij}}{\partial \varepsilon_{kl}} \text{ (elasto-optic)}$$

The second is a **third** derivative of the energy, and it comes from one `jvp` of the second-order energy at frozen first-order wavefunctions — the $2n+1$ theorem again. Through the thermodynamic identity of Tanner, Bousquet and Janolin it gives the four electrostriction tensors m , q , M and Q .

```
from pypresso.response import elastic_constants, electrostriction

es = electrostriction(calculation, scf)      # solves the strain response itself
print(es.m, es.q, es.M, es.Q)              # the four electrostriction tensors
print(es.photoelastic)                     # the elasto-optic tensor
print(es.elastic.voigt)                    # C_ij in Voigt notation, GPa

# or the elastic constants alone, from a strain response you already have
C = elastic_constants(calculation, scf.wavefunctions, scf.eigenvalues,
                      scf.density, es.strain)
print(C.tensor.shape, C.voigt)
```

Converged silicon gives $C_{11} = 198.5$ GPa and $C_{12} = 68.9$, against a measured 165.7 and 63.9; the three independent components of $\partial\epsilon/\partial\varepsilon$ match a central difference over re-converged strained cells to **2e-4**, which is that difference’s own floor, and the rank-4 tensor is cubic to 3e-14 with nothing imposing it.

Refuses. Both are norm-conserving, `nspin = 1`, insulators and **clamped-ion**. A diverged first-order solution is refused rather than consumed in silence (`require_converged_responses`) — at QE’s `alpha_mix = 0.7` the strain response of a *slab* diverges, and consuming it gave a C_{ijkl} that was not even symmetric under $C_{ijkl} = C_{klij}$. The elastic constants additionally need the whole k -grid rather than a wedge, because their functional builds its own density and symmetrises it as a scalar inside the chain rule.

7.1 The piezoelectric tensor

The polarization a strain induces is the stress an electric field induces — one mixed second derivative read either way round:

$$e_{k,ij} = \frac{\partial P_k}{\partial \varepsilon_{ij}} = \frac{\partial \sigma_{ij}}{\partial \mathcal{E}_k} = -\frac{1}{\Omega} \frac{\partial^2 E}{\partial \varepsilon_{ij} \partial \mathcal{E}_k}$$

It is the Born charge’s construction with one coordinate changed. Z^* is a jvp of the *force* along the field’s response; the force is `jax.grad` of the frozen-state energy in the atomic positions and the stress is `jax.grad` of the same functional in a strain, so this is a jvp of the *stress* along the same response — three of them, one per field direction, on top of a dielectric constant that was going to be solved anyway. The orthonormality constraint contributes nothing here, because $\langle \psi | \hat{S} | \psi \rangle$ is a sum over a sphere of integers and carries no cell.

```
calc = Calculator.from_file("alas.in", pseudo_dir="pseudo")

piezo = calc.get_piezoelectric_tensor()
print(piezo.voigt)           # (3, 6) in C/m^2, columns xx yy zz yz xz xy
print(piezo.e14)            # -0.7638: zincblende's one component
print(piezo.e.shape)        # (3, 3, 3) in e/bohr^2, e[k, i, j]
print(piezo.dielectric.epsilon) # the field response is shared, not re-solved
```

`pw.x` computes no piezoelectric tensor — the word occurs once in the whole distribution, in a citation in a comment — and Elk’s task 380 reaches it by a finite difference of the Berry-phase polarization over one full ground state per strain tensor. So the validation is internal, and there are four independent statements. Silicon is centrosymmetric and its whole tensor vanishes (2.4e-5 C/m² against AlAs’s 0.764 from the same code, with nothing imposing it on a `nosym` run); AlAs is $\bar{4}3m$ and only $e_{14} = e_{25} = e_{36}$ survive, to 1.7e-14; an eight-point wedge reproduces the sixty-four-point closed grid to **4.5e-9**; and the same mixed derivative contracted the other two ways — `zstar_eu.f90`’s expression with a strain label, and the strain response against the field’s bare perturbation — agree to **6.2e-15** and **1.3e-7**. What anchors the sign and the normalisation is that the *same* assembly run in the position coordinate is the Born charge, and that is `ph.x`’s number.

Refuses. Clamped-ion — the atoms are carried by the strain and do not relax, which is also what Elk’s task computes; the internal-strain term that makes the constant comparable with experiment is not here, and for zincblende it very nearly cancels this one. And **polar crystals by name**: what the derivative above gives is the *improper* tensor, and the proper piezoelectric response differs from it by $\delta_{ki}P_j - \delta_{ij}P_k$, which needs the spontaneous polarization itself. Those terms vanish identically whenever the two Cartesian labels they pair are different — so e_{14} of a zincblende crystal never carries the ambiguity — and they vanish for every component of a class that admits no invariant vector, which is what is checked. **Ultrasoft and PAW**, and that one is a gap rather than a missing term: nothing in the assembly is norm-conserving and the displacement leg of it is the Born charge, which is validated on all three kinds, but every ultrasoft and PAW crystal here is centrosymmetric and agrees with zero whatever is wrong. Beyond that: insulators, `nspin = 1`, and everything the Sternheimer regime and the differentiable cell refuse (a metal, a spin spiral, a magnetic field, a shifted grid run with `nosym`).

7.2 Spin-polarized response

The Sternheimer equation is solved per spin channel,

$$(H_\sigma - \varepsilon_{n\sigma}S + \alpha_{pv}Q_\sigma) |\Delta\psi_{n\sigma}\rangle = -P_{c,\sigma}^+ \Delta V_\sigma |\psi_{n\sigma}\rangle,$$

with the occupied manifold $P_{c,\sigma}$ built from *that channel's* filling: $N_{\uparrow}^{\text{occ}} = \text{NINT}(n_{\uparrow})$ and $N_{\downarrow}^{\text{occ}} = \text{NINT}(n_{\downarrow})$, which for a magnetic insulator are different numbers. Quantum ESPRESSO never meets this because LSDA doubles nks there and each spin channel is a separate k -point.

```
# chi_0 for a spin-polarized run -- the part that is validated unconditionally.
from pypresso.response.sternheimer import make_sternheimer

solver = make_sternheimer(calculation, result, spin_polarized=True)
drho = solver.chi0(dv)          # dv is (2, n1, n2, n3) on the dense grid
```

Validated against a central difference of the density under a spin-dependent probe potential — 1.8×10^{-6} on an antiferromagnetic hydrogen chain (a smeared metal) and 1.1×10^{-6} on triplet O₂ (the sliced branch, ultrasoft) — and against the identity that a cell with no magnetization gives the same ε_∞ at `nspin = 1` and `nspin = 2`: 6.2×10^{-14} on 13.806646105. The probe must differ between the channels, because χ_0 is block diagonal in spin and one equal in both would not tell the blocks apart.

Refuses. `tot_magnetization` together with a smearing: the two channels have separate Fermi levels and `Smearing` carries one scalar `ef`. A filling whose boundary lands inside a *degenerate multiplet* — a Hund's-rule atom is exactly this case — where the solve converges and returns an answer 100 % away from a finite difference, because which member falls below the cut is arbitrary; that is the physics, not a missing term. Born effective charges, the dynamical matrix, the strain response, the elastic constants, the Raman tensor and the electrostriction tensors for `nspin = 2`: the solve is spin-polarized, their second-derivative assembly is not. The *screened* response of a magnetic system with vacuum: for `nspin = 2` the kernel `dv_of_drho` is the second derivative of the LSDA energy in the two channel densities and diverges wherever a channel density reaches zero — measured on triplet O₂, 1504 of 91125 grid points and exactly 1504 NaN. And any potential-only meta-GGA (tb09, bj06), which has no total energy to differentiate — refused by name now rather than surfacing as `v_of_rho` asking for `tau`.

8 Optical spectra and excitons: TDDFT

The response above is a *static* one, solved orbital by orbital. An absorption spectrum needs the frequency axis and needs the susceptibility as a **matrix** in reciprocal lattice vectors, so this is the one place where a sum over states earns its keep. `run_absorption` builds

$$\chi_{\mathbf{G}\mathbf{G}'}^0(\omega) = \frac{1}{\Omega} \sum_{\mathbf{k}} \sum_{ij} \frac{(f_i - f_j) \rho_{\mathbf{G}}^{ij} \rho_{\mathbf{G}'}^{ij*}}{\varepsilon_i - \varepsilon_j + \omega + i\eta}, \quad \rho_{\mathbf{G}}^{ij} = \langle u_i | e^{-i\mathbf{G}\cdot\mathbf{r}} | u_j \rangle$$

and solves the Dyson equation with an exchange-correlation kernel:

$$\epsilon^{-1} = 1 + X(1 - X - \tilde{f}_{\text{xc}}X)^{-1}, \quad X = v^{1/2}\chi^0v^{1/2}, \quad \tilde{f}_{\text{xc}} = v^{-1/2}f_{\text{xc}}v^{-1/2}$$

The **bootstrap** kernel of Sharma, Dewhurst, Sanna and Gross (*PRL* **107**, 186401 (2011); Elk's `fxctype = 210`) is that equation's other half, so the two are solved together to self-consistency:

$$f_{xc}^{BS} = -\frac{\epsilon^{-1}(\mathbf{q}, 0) v(\mathbf{q})}{\epsilon_0(\mathbf{q}, 0) - 1}, \quad \epsilon_0 = 1 - v\chi^0$$

It is parameter-free and diverges as $1/q^2$, which is what binds an electron-hole pair; ALDA's kernel is finite at $\mathbf{q} = 0$ where v diverges, so its head and wings vanish *identically* and no amount of it can bind one.

Quantum ESPRESSO has no counterpart. TDDFPT/ is a Liouville-Lanczos solver with RPA and ALDA; it has no bootstrap kernel and no Dyson equation in $\mathbf{G}$ space. The reference is Elk.

```
from pypresso.workflows import run_absorption
import numpy as np

omega = np.arange(0.0, 0.6, 0.005)          # Ry
spectrum = run_absorption(system, pseudos, scf.density, omega,
                          kernel="bootstrap", # or rpa / alda / lrc / bootstrap-1
                          nbnd=60, ecut_response=8.0,
                          broadening=0.012, scissor=0.05)

print(spectrum.absorption)                  # Im eps_M, isotropic average
print(spectrum.frequencies_ev)              # the same grid in eV
print(spectrum.alpha, spectrum.iterations)  # the LRC-equivalent head, and the
                                           # bootstrap's fixed-point count
print(spectrum.static_residual)             # how far the band truncation reaches
```

`chi_0` is the whole cost and the kernel is nearly free, so comparing kernels means building it once: `independent_response` returns it and `solve_dyson` screens it, which is what `run_absorption` does internally and what the notebook shows.

Three knobs here are this code's own. `ecut_response` (Elk's `gmaxrf`) selects the $\mathbf{G}$ -set the matrix is built over and has its own convergence: too small and the local-field effect is missing, and the cost is quadratic in it. `scissor` is *part of the method* rather than a convenience — the paper's Eq. (3) replaces χ^0 by a gap-corrected model response, and the velocity matrix elements are renormalised by $e_{ij}/(e_{ij} \mp \Delta)$ along with it. And `static_residual` is the diagnostic for the one error this phase cannot refuse: how many empty states the sum ran over. It is $\epsilon_M(0)$ from this run, in RPA, minus the same quantity from the Sternheimer solve, which is band-complete; it is kernel-matched on purpose, since differencing a bootstrap number against a Sternheimer one that contains f_{xc} would measure the kernel instead.

A trap that is invisible in the answer: the macroscopic dielectric function is the inverse of the 3×3 *head* of ϵ^{-1} , not the head of the inverse of the whole matrix. The second is the microscopic ϵ , whose head in RPA is exactly the no-local-field result; it is smooth, positive, has the right peaks and is 9% too large on silicon. Both are returned (`epsilon` and `epsilon_no_local_fields`), because the gap between them is the local-field effect.

Validated by an identity rather than against another code's spectrum: the same $\epsilon_M(0)$ is reached by this sum over states plus a Dyson inversion and by the projected conjugate-gradient solve of `dielectric_tensor`, which shares no machinery with it and never sees an empty state. On silicon the two agree to **1.3e-2** on a constant of 22 — and that residue is the band truncation, which is why it is reported rather than tuned away. The screening = "hartree" switch on `dielectric_tensor` exists for this comparison: the identity holds only

when both sides use the same kernel.

Refuses. Finite $\mathbf{q}$ (the perturbed states would live at $\mathbf{k} + \mathbf{q}$); ultrasoft and PAW (every matrix element gains the augmentation charge $Q_{ij}(\mathbf{G})$); metals (the $f_i - f_j$ weight kills the $i = j$ term, so the intraband Drude response is absent entirely — Elk’s `tdfft1r` does not add it either); `nspin` $\neq 1$, noncollinear magnetism, spin-orbit coupling, spin spirals and a Hubbard U ; a symmetry-reduced k -set, because symmetrising $\chi_{\mathbf{G}\mathbf{G}'}^0$ rotates *two* $\mathbf{G}$ indices at once and the rank- n symmetriser is Cartesian; the `alda` kernel with a gradient-corrected functional, since ALDA’s kernel is pointwise in the density and a GGA’s is not; and **non-convergence of the bootstrap fixed point is an error**, as it is in `tdfft1r.f90`, because a spectrum from a kernel still in motion is the spectrum of nothing.

8.1 The conductivity tensor, the Kerr angle and the anomalous Hall effect

The whole complex tensor, interband plus a Drude term:

$$\sigma_{ij}(\omega) = \frac{i}{\Omega} \sum_{\mathbf{k}} \sum_{n,m} \frac{W_n (1 - f_m)}{\epsilon_{mn}} \left[\frac{v_{nm}^i v_{mn}^j}{\omega - \epsilon_{mn} + i\eta} + \frac{\overline{v_{nm}^i v_{mn}^j}}{\omega + \epsilon_{mn} + i\eta} \right] + \frac{\omega_{p,ij}^2}{4\pi(\gamma - i\omega)}$$

with $v = \partial H / \partial \mathbf{k}$ from one `jvp` of $H(\mathbf{k})$ at a frozen sphere. *That* is the one thing here not transcribed from `dielectric.f90`, and it is not a refinement: `epsilon.x` accumulates $\langle \psi_1 | \mathbf{G} | \psi_2 \rangle$, and $[H, \mathbf{r}] \neq \mathbf{p}$ when the pseudopotential is nonlocal.

What is new beside `run_absorption` is the **antisymmetric** part, σ_{xy} — the response that rotates light reflected off a magnet, and whose static limit is the intrinsic anomalous Hall conductivity. Time reversal forces it to zero, so a nonmagnetic crystal has none however strong its spin-orbit coupling; a *collinear or coplanar* magnet without spin-orbit coupling has none either, because a spin rotation composed with conjugation survives as an antiunitary symmetry. (A *noncoplanar* magnet does have one with no spin-orbit coupling at all — the topological Hall effect — so the usual pairing is a route and not a theorem.)

```
calc = Calculator.from_file("ni-soc.in", pseudo_dir="pseudo")

sigma = calc.get_optical_conductivity(nbnd=36, window=0.6, nw=200)
print(sigma.sigma_s_per_cm.shape)    # (nw, 3, 3) in S/cm
print(sigma.kerr)                    # complex Kerr angle, degrees, per frequency
print(sigma.hall_conductivity)       # the static antisymmetric part, S/cm
print(sigma.plasma_ev)               # hbar omega_p, eV, a (3, 3) tensor
print(sigma.dielectric)              # 1 + 4 pi i sigma / omega
print(sigma.fsum, sigma.band_cut_gap) # the two diagnostics to read
```

There is no reference output for any of it, so the validation is internal and analytic. The *symmetric* part reproduces the independent-particle dielectric function of the TDDFT stack to **3.6e-14**, which is a different assembly (a Dyson solve over a response sphere) reaching `ph.x` through `dielectric_tensor`. The **f-sum rule** is the absolute anchor: $\int_0^\infty \text{Re } \sigma_{aa} d\omega$ must equal $(\pi/2\Omega) \sum W_n \langle n | \partial^2 H / \partial k_a^2 | n \rangle$, and the familiar $\pi n_e / 2$ is only the *local* Hamiltonian’s special case — silicon’s diamagnetic weight is measurably **0.943** and not one. Measured against a central difference of the velocity operator, the ratio is 1.258 on a 4^3 grid, 1.045 on 6^3 and **0.990** on 8^3 : it converges in the *k-grid*, because what separates the two is $\sum_k w_k \partial^2 \epsilon_n / \partial k^2$, which vanishes over the zone by periodicity and not on

a coarse mesh. Aluminium's plasma frequency comes out at **12.98 eV** against a free-electron 16.27, which is the zone-boundary gaps removing Fermi surface, and its tensor is isotropic to **1.7e-4 eV** with nothing imposing that.

Refuses. Ultrasoft and PAW, for the Kubo curvature's reason: with a moving $\hat{S}$ the current operator carries $\epsilon_n \partial S / \partial k$ off the diagonal, and that term is identically zero for a norm-conserving dataset, so nothing validated here can see whether its convention is right. A **spin spiral**, whose two spinor components live on different spheres. A **symmetry-reduced k-set**, because the antisymmetric part is an axial vector — the escape is the whole unshifted grid, which is closed under the point group. And **nspin = 2**, whose two channels are two band structures whose conductivities add; a spinor run is the regime a magneto-optical spectrum wants anyway, since σ_{xy} needs spin-orbit coupling and collinear spin has none.

Two things are reported rather than refused, and both must be read. `band_cut_gap` is the gap between the last band summed and the first dropped: stopping *inside* a degenerate multiplet keeps some members and drops others, and silicon's antisymmetric σ — zero by time reversal — comes out at 4.0e-13 when the cut falls at a real gap and at 1.0e-5 when it does not. And the **anomalous Hall conductivity of a metal is mesh-hungry**: the static Berry-curvature integral is concentrated on near-degeneracies at the Fermi surface, so it is quoted with its k -convergence and not on its own.

9 Topological invariants

Everything is built from one primitive — the overlap of occupied manifolds at neighbouring k -points,

$$M_{mn}^{(\mathbf{b})} = \langle u_{m\mathbf{k}} | \hat{S} | u_{n,\mathbf{k}+\mathbf{b}} \rangle$$

because a determinant of overlaps is blind to the unitary mixing a degenerate eigensolver leaves. The Berry curvature is the phase of a plaquette of them and the Chern number their lattice sum, which is an **exact integer on any mesh**:

$$F_{12}(\mathbf{k}) = \ln [U_1(\mathbf{k})U_2(\mathbf{k}+\hat{1})U_1(\mathbf{k}+\hat{2})^{-1}U_2(\mathbf{k})^{-1}], \quad C = \frac{1}{2\pi i} \sum_{\mathbf{k}} F_{12}(\mathbf{k})$$

Z_2 has two independent routes — the flow of Wannier charge centres, and, when there is an inversion centre, the parity product over the eight TRIM points:

$$(-1)^\nu = \prod_{i=1}^8 \prod_{n \text{ occ}}^{\text{Kramers}} \xi_{2n}(\Gamma_i)$$

```
from pypresso.workflows.topology import (run_z2, run_z2_3d,
                                         run_berry_curvature)

z2 = run_z2(system, pseudos, scf.density, axis=2) # Wilson loop by default
print(z2.z2, z2.gap_step) # read gap_step before believing the integer

chern = run_berry_curvature(system, pseudos, scf.density, shape=(12, 12))
print(chern) # an exact integer on any mesh

nu = run_z2_3d(system, pseudos, scf.density) # the four 3D indices
print(nu) # (nu0; nu1 nu2 nu3)
```

Refuses. A collinear spin-polarized run (two Hamiltonians, so two independent sets of invariants and no statement of which is meant); a Z_2 without spin-orbit coupling, where every invariant is zero for a reason that has nothing to do with the band structure; the parity route without an inversion centre; and a manifold that is not separated from the band above it (`gap_tol`). A PAW run needs the converged `becsum` passed beside the density — `run_z2(..., becsum = scf.becsum)`, which `Calculator` does for you — because `ddd_paw` is a property of the wavefunctions and cannot be rebuilt from a density; leaving it out is worth 1.03 eV in the eigenvalues, and an invariant computed from those is still an integer.

Running both Z_2 routes is the check. Where they disagree the parity one is the answer, because it has no mesh and the Wilson route does. **Two things bite in a plane-wave code and both are silent:** neighbouring k -points do not share a G -sphere, so coefficients are aligned by Miller index; and the wrap at the zone edge is a *shift* of that index, without which the Chern number comes out smooth and non-integer.

9.1 The curvature as a map: the Kubo route

Two constructions of the same quantity, chosen by name:

$$\Omega_n^{12}(\mathbf{k}) = -2 \operatorname{Im} \sum_{m \neq n} \frac{A_{nm}^1 A_{mn}^2}{(\varepsilon_n - \varepsilon_m)^2}, \quad A_{nm}^a = \langle \psi_n | \partial_{k_a} H - \varepsilon_n \partial_{k_a} S | \psi_m \rangle$$

with `method="kubo"`, against the lattice field strength of Fukui, Hatsugai and Suzuki with `method="fhs"` (the default). The velocity operator is one `jvp` of $H(\mathbf{k})$ at a frozen plane-wave sphere — no dense $H(\mathbf{k})$ is ever formed — and the tangent for a crystal direction d is the reciprocal lattice vector $\mathbf{b}_d$.

```
from pypresso.workflows.topology import run_berry_curvature

curvature = run_berry_curvature(
    system, pseudos, density, shape=(4, 4), nocc=4, nbnd=12, method="kubo"
)
curvature.curvature           # (4, 4), Omega(k)
curvature.curvature_by_band   # (4, 4, nocc) -- see the refusal box
curvature.truncation          # what the highest empty band was worth
```

Validated against the FHS flux, which shares no machinery with it: a centred plaquette shrunk around one k -point of zincblende AlAs converges onto the pointwise Kubo value to 3.2×10^{-3} , and the whole 24×24 mesh agrees plaquette by plaquette to 1.45×10^{-4} , improving 65-fold over a sixfold refinement. The velocity matrix elements themselves reproduce a central finite difference of $H(\mathbf{k})$ to 1.8×10^{-9} , and on centrosymmetric silicon the curvature vanishes *pointwise* to 3.5×10^{-5} .

Refuses. `method="kubo"` is for the map, never for the integer: its denominator $(\varepsilon_n - \varepsilon_m)^{-2}$ is singular at a degeneracy and its Brillouin-zone sum is an ordinary Riemann sum, which converges to an integer without ever being one — `method="fhs"` is the default and is exact on any mesh. **Ultrasoft and PAW are refused by name:** the $\varepsilon_n \partial_k S$ term is identically zero for a norm-conserving dataset, so no validation here can see whether its off-diagonal convention is right, and $q_{ij}(\mathbf{b})$ is the second missing piece; `method="fhs"` carries both correctly on all three dataset kinds. The sum over empty states is truncated at `nbnd` and the truncation is *reported* rather than hidden — read

BerryCurvature.truncation, or truncation_abs wherever symmetry forces the curvature to vanish, since the ratio is then noise over noise. curvature_by_band is gauge invariant only for a non-degenerate band; inside a degenerate multiplet only the sum over the members is defined, and curvature (the manifold total) always is.

10 Effective masses and site angular momenta

Two quantities Elk computes and pw.x does not, both at the cost of one NSCF at a single k-point.

10.1 The effective mass tensor

$$\left(\frac{1}{m^*}\right)_{ab} = \frac{1}{2} \frac{\partial^2 \varepsilon_n(\mathbf{k})}{\partial k_a \partial k_b},$$

in units of $1/m_e$. The factor of a half is Rydberg atomic units and not a normalisation: $\hbar^2/2m_e$ is exactly 1 Ry bohr², so a free electron has $\varepsilon = |\mathbf{k}|^2$ and the tensor is the identity.

The first derivative is analytic and only the second is a difference. $\partial \varepsilon_n / \partial k$ is the generalised Hellmann–Feynman expression built from one jvp of $H(\mathbf{k})$; its second derivative is not reachable the same way, because an individual band’s $|\partial_k \psi_n\rangle$ is not what the Sternheimer projector produces (it removes the whole occupied manifold, and the band whose mass is wanted is usually empty). So the construction is one central difference of an *exact* first derivative — six stencil points against Elk’s twenty-seven, and no polynomial fit. Elk’s route, differencing eigenvalues, is kept beside it as method="eigenvalue" because it shares nothing with the velocity operator.

```
from pypresso import Calculator

calc = Calculator.from_file("si.scf.in")
mass = calc.get_effective_mass((0.0, 0.0, 0.0), nbnd=10) # crystal coordinates

mass.inverse_mass          # (nbnd, 3, 3) in 1/m_e, cartesian
mass.principal(band=7)     # the three principal masses and their axes
mass.density_of_states_mass(band=7) # (m_x m_y m_z)^(1/3)
mass.truncation            # the O(h^2) the Richardson step removed
mass.multiplets            # per multiplet, including the degenerate ones
```

On two-atom silicon at Γ against the vendored **Elk** binary (all-electron LAPW, PBE, $a = 10.26$ bohr) with a PAW dataset here: the non-degenerate Γ_{1v} agrees to **0.02%** (0.8600902 against 0.8603044, i.e. $m^* = 1.16267 m_e$ against 1.16238) and $\Gamma_{2'c}$ to 0.36% ($m^* = 0.170439 m_e$). The two routes here agree with each other to 1.2×10^{-5} , and the tensors come out isotropic to 2.9×10^{-8} with nothing imposing cubic symmetry.

Refuses. A band inside a **degenerate multiplet** has no tensor of its own — the eigensolver’s arbitrary rotation within the manifold rotates it — so those bands come back as nan and what is reported instead is the multiplet’s *summed* inverse mass, which is the trace and is invariant. Silicon’s threefold $\Gamma_{25'}$ is the case, and Elk’s own output is the evidence: it prints $-5.8938, -3.7690, -3.7690$ for three states whose energies agree to 5×10^{-8} Ha. A **spin spiral** is refused, since a k-stencil moves both spheres. delta is the stencil’s outermost displacement in both routes; the $O(h^2)$ truncation is removed

by a Richardson step and then *reported* in truncation rather than tuned away.

10.2 $\langle L \rangle$, $\langle S \rangle$ and $\langle J \rangle$ on each atom

Where the orbital moment of a spin-orbit magnet actually sits. From one site density matrix per atom and shell, built out of the same projection a projected density of states uses,

$$\rho_{(ms),(m's')}^a = \sum_{n\mathbf{k}} w_{n\mathbf{k}} c_{ms,n\mathbf{k}}^a \overline{c_{m's',n\mathbf{k}}^a}, \quad c = \langle \phi | S | \psi \rangle,$$

and then $\langle L_i \rangle = \sum_s \text{Tr}_m(L_i \rho_{ss})$ and $\langle S_i \rangle = \frac{1}{2} \sum_m \text{Tr}_s(\sigma_i \rho_{mm})$, with $J = L + S$. The L matrices are the complex-basis ones conjugated into the *real* spherical harmonics the code works in, using the same `rot_ylm` unitary the spin-orbit projectors are built from.

```
from pypresso import Calculator

calc = Calculator.from_file("ni.scf.in") # noncolin, lspinorb, nosym
momenta = calc.get_angular_momenta()

print(momenta.table()) # what Elk writes to LSJ.OUT
momenta.atoms[0].l      # <L> on atom 1, cartesian, units of hbar
momenta.atoms[0].j      # <J> = <L> + <S>
momenta.total_orbital   # summed over the cell
```

Validated by identities, not by another code's number — Elk integrates over a muffin tin of a stated radius and this projects onto an orbital set, so the two differ by definition the way Löwdin and muffin-tin charges do. $\langle L \rangle$ is *quenched to zero* without spin-orbit coupling, measured at 1.7×10^{-16} on silicon; with the coupling on, fully-relativistic nickel gives $\langle L_z \rangle = 0.0364767 \hbar$ against a measured orbital moment of about $0.05 \mu_B$, with $|L|/|S| = 0.11665$ against an experimental $m_L/m_S \approx 0.1$ and $\langle L \rangle \parallel \langle S \rangle$ (Hund's third rule). Driving the moment along z , x and y gives $|\langle L \rangle| = 0.0364767$ in all three, with a spread of $7.3\text{e-}11$ over the three cubic axes, with nothing imposing that a magnitude is a scalar.

Refuses. A **symmetry-reduced k-set**: $\langle L \rangle$ and $\langle S \rangle$ are vectors, and summing a wedge sums a wedge — the axial-vector group average with $\det(R)$ and the time-reversal sign is not written here. Run the whole grid with `nosym = .true.`, unshifted so that it is closed under the point group. A **fully-relativistic ultrasoft or PAW** dataset, because the spinor overlap's off-diagonal spin blocks are `qq_so` and the projection here applies the scalar S to each component; a fully-relativistic *norm-conserving* dataset has $S = 1$ and is exact. And a **spin spiral**, whose two components do not share a set of atomic orbitals. `pw.x`'s `lorbm` computes the *cell's* orbital magnetization and has no per-atom counterpart to check this against.

11 Things with no `pw.x` counterpart

Spin spirals by the generalized Bloch theorem: a spiral of wavevector $\mathbf{q}$ is a spinor whose two components live on *different* plane-wave spheres,

$$\psi_{\mathbf{k}}(\mathbf{r}) = \begin{pmatrix} e^{i(\mathbf{k}-\mathbf{q}/2)\cdot\mathbf{r}} u^\uparrow(\mathbf{r}) \\ e^{i(\mathbf{k}+\mathbf{q}/2)\cdot\mathbf{r}} u^\downarrow(\mathbf{r}) \end{pmatrix}$$

so in the rotated frame the density is lattice periodic and the SCF, the functional and the mixer are untouched. Only two terms carry $\mathbf{q} = |\mathbf{k} \mp \mathbf{q}/2 + \mathbf{G}|^2$ and $v_{\text{nl}} =$ and $dE/d\mathbf{q}$ at frozen coefficients is `jax.grad` of them.

```
from pypresso.workflows.spiral import (run_spiral_scan, relax_spiral_q,
                                       heisenberg_exchange)

scan = run_spiral_scan(system, pseudos, wavevectors) # E(q)
best = relax_spiral_q(system, pseudos) # BFGS on the reciprocal metric
print(best.wavevector, best.scf.total_energy)

J = heisenberg_exchange(scan, shells) # fit E(q) for the exchange constants
```

`heisenberg_exchange` fits $E(\mathbf{q}) - E(0) = m^2 \sum_{\mathbf{R}} J(\mathbf{R}) [1 - \cos(\mathbf{q} \cdot \mathbf{R})]$ over a set of lattice shells, which is how a spiral scan becomes a spin model. `compute_spiral_gradient` gives $dE/d\mathbf{q}$ on its own.

Also here: **per-atom magnetic fields**, the **topological invariants** of the previous section, and **rank- n tensor symmetrisation** — `symme.f90`'s `symmatrix3/symtensor3` written at any rank.

Refuses. A spiral with spin-orbit coupling, permanently — it breaks the theorem, and Elk refuses it too; a spiral with symmetry, until the spin space group is written, so a spiral needs `nosym`; a spiral with ultrasoft/PAW; and spiral relaxation under a magnetic field, whose energy is outside the reported total, so the state is stationary for a different functional than the one being differentiated.

12 Performance: dials, memory and GPUs

Two batching dials control how much is in flight, and both defaults follow the platform: on a CPU, QE's loop — one k -point and one band at a time, which is what a *cache* wants — and on a GPU the whole of both axes, which is what a card wants. Either end is reachable everywhere:

```
PYPRESSO_K_BATCH=all PYPRESSO_BAND_BATCH=all python3 run.py
```

or per call, `run_scf(..., k_batch=...)`. The chunk size changes the order contributions are summed in and nothing else — round-off, $\sim 1\text{e-}15$ Ry.

Why the default is per platform. A GPU has no cache to fit in and inverts the conclusion: on a ten-atom metal, `k=1,b=1` costs **4.5x** what `k=all,b=all` does on the same card, sixteen-atom silicon runs at **0.20x** — an outright loss against four CPU cores — and at thirty-two atoms `b=1` did not finish at all. So a run on a device gets both batched unless it says otherwise, and the plausible half-measure `k=all,b=1` is never the default: it is the *worst* setting available, because it buys the batched mode's memory with the looped mode's kernel launches. An explicit argument beats the environment variables, which beat the platform.

Measured GPU speedups (one H200 against one CPU core, per SCF iteration): 2x on a cheap metal, 3-5x on norm-conserving silicon, 9-21x with an augmentation charge or a magnetization, and **83-115x on spin-orbit**. Energies agree to $\sim 1\text{e-}9$ Ry — the same agreements the CPU already has.

Memory. Spin-orbit is where it bites: a 20-atom bismuth cell with a relativistic ultrasoft dataset peaks at **34.7 GB**, which fits a 141 GB card and does not fit a 32 GB one. Everything else in that sweep sits under 2 GB.

A response costs memory according to how it is differentiated, which is worth knowing before starting one. A *forward* response — the Sternheimer solves behind ϵ and Z^* — costs 0.7–1.0 GB over the SCF that produced the state, and a *forward-over-reverse* one — the dynamical matrix, the Raman tensors — 1–2 GB. The *stress* is the exception and it is a *reverse* derivative through the radial transforms: **10.5 GB** on eight-atom ultrasoft silicon, ten times any of the others, and an order more than the SCF it came from. Measure your own with `tools/gpu/phase5.py <case> --property <name>`.

Precision. Every dtype comes from `config.Precision`; x64 is enabled before any array exists, and all validation is float64. float32 is a performance mode and **never one a correctness claim is made in**.

13 Accuracy summary

Against Quantum ESPRESSO on the same input:

quantity	agreement
total energy, norm-conserving LDA	$\sim 1\text{e-}9$ Ry
ultrasoft, PAW	$\leq 3\text{e-}9$ Ry
PBE / revPBE / PBEsol	$\leq 6\text{e-}9$ Ry
metals, every smearing and tetrahedra	$2.5\text{e-}8$ Ry
collinear spin	$1.2\text{e-}9$ Ry
noncollinear magnetism	$2.8\text{e-}9$ Ry
spin-orbit coupling	$\leq 1.3\text{e-}8$ Ry
DFT+U	$\leq 6.7\text{e-}9$ Ry
band structure	0.0002 eV
forces	$\leq 2\text{e-}5$ Ry/bohr
stress (13 cases)	$\leq 2.7\text{e-}7$ Ry/bohr ³
relaxed geometry	$1\text{e-}6$ bohr
dielectric constant (NC, US, PAW)	$\leq 1.2\text{e-}4$
Born effective charges (NC)	every printed digit
phonons at Γ , silicon	0.05 cm^{-1}
phonons at Γ , silicon (US, PAW)	$0.02\text{--}0.03\text{ cm}^{-1}$
phonons at Γ , aluminium (metal)	0.002 cm^{-1}
Raman/IR spectra vs <code>dynmat.x</code>	every printed digit

One case does not close, and is recorded rather than explained. `bi10-soc` — ten bismuth atoms with a fully-relativistic dataset — sits **1.9e-4 Ry** from QE on a total of -1477.737 , which is $1.3\text{e-}7$ relative. Both codes choose the same symmetry operations, the same k -point, the same grid and the same G -vectors, and both converge. It is in the test suite at that tolerance, not hidden.

Where to go next

- `README.md` — the short tour and a first calculation
- `notebooks/` — 26 worked examples on concrete systems, executed and committed, each with a `.md` export beside it

- `PERFORMANCE.md` — the running performance log, CPU and GPU
- `PLAN.md` — the architecture, the phase breakdown, the trap catalogue
- `GPU.md` — the GPU roadmap