

# defumat

what it computes, and how to ask for it

Plane-wave DFT in Python and JAX, validated against Quantum ESPRESSO

September 2, 2026

## Abstract

User documentation for the feature set: every capability, the equation behind it, a snippet that runs it, what it was validated against, and *what it refuses*. Two conventions run through the whole document.

**Everything is validated against Quantum ESPRESSO on the same input**, and the agreement is quoted per feature rather than claimed in general. Where a number is missing it is because the comparison does not exist, and that is said rather than hidden.

**Anything not implemented is refused by name**, with an error saying what is missing — never silently approximated. The “Refuses” notes are as much of the documentation as the features: they are the promise that a run which starts is a run whose physics is there.

---

**One equation is written down and the rest are derived from it.** Almost everything below is a derivative of the total energy, taken by the compiler rather than by hand:

$$E[\{\psi\}, \tau, \varepsilon] = T_s + E_H + E_{xc} + E_{\text{loc}} + E_{\text{nl}} + E_{\text{Ewald}}$$

| quantity                                 | what it is                                                      | how it is obtained                        |
|------------------------------------------|-----------------------------------------------------------------|-------------------------------------------|
| $v_{xc}(\mathbf{r})$                     | $\delta E_{xc} / \delta n$                                      | jax.grad of the energy density            |
| $\mathbf{F}_I$                           | $-\partial E / \partial \boldsymbol{\tau}_I$                    | jax.grad at frozen $\psi$                 |
| $\sigma_{ij}$                            | $-\Omega^{-1} \partial E / \partial \varepsilon_{ij}$           | jax.grad along a strain                   |
| $C_{I\alpha, J\beta}$                    | $\partial^2 E / \partial \tau_{I\alpha} \partial \tau_{J\beta}$ | jvp of the force                          |
| $Z_{I, \alpha\beta}^*$                   | $\partial F_{I\alpha} / \partial \mathcal{E}_\beta$             | jvp of the force along the field response |
| $\partial \epsilon_{ij} / \partial \tau$ | Raman                                                           | jvp of the second-order energy            |

The “frozen  $\psi$ ” is what makes this exact rather than approximate: at self-consistency the energy is stationary with respect to the wavefunctions, so the terms from  $\partial \psi / \partial \lambda$  vanish. That is the Hellmann-Feynman argument, and taking it one order up is the  $2n+1$  theorem, which is why a *third* derivative needs only first-order wavefunctions.

# Contents

|           |                                                                                          |           |
|-----------|------------------------------------------------------------------------------------------|-----------|
| <b>1</b>  | <b>Input and invocation</b>                                                              | <b>3</b>  |
| <b>2</b>  | <b>Ground state</b>                                                                      | <b>5</b>  |
| <b>3</b>  | <b>Bands, densities of states, projections</b>                                           | <b>6</b>  |
| <b>4</b>  | <b>Pseudopotentials and functionals</b>                                                  | <b>6</b>  |
| <b>5</b>  | <b>Spin, magnetism and spin-orbit</b>                                                    | <b>7</b>  |
| <b>6</b>  | <b>Forces, stress and relaxation</b>                                                     | <b>10</b> |
| <b>7</b>  | <b>Response: dielectric, phonons, Raman</b>                                              | <b>11</b> |
| 7.1       | The piezoelectric tensor . . . . .                                                       | 15        |
| 7.2       | Spin-polarized response . . . . .                                                        | 16        |
| <b>8</b>  | <b>Optical spectra and excitons: TDDFT</b>                                               | <b>16</b> |
| 8.1       | The conductivity tensor, the Kerr angle and the anomalous Hall effect . . .              | 18        |
| 8.2       | The shift current: the bulk photovoltaic effect . . . . .                                | 19        |
| 8.3       | Second-harmonic generation . . . . .                                                     | 21        |
| <b>9</b>  | <b>Topological invariants</b>                                                            | <b>22</b> |
| 9.1       | The curvature as a map: the Kubo route . . . . .                                         | 23        |
| 9.2       | The Berry-phase polarization . . . . .                                                   | 24        |
| 9.3       | The magnetoelectric tensor . . . . .                                                     | 25        |
| <b>10</b> | <b>Effective masses, site angular momenta, nesting and structure factors</b>             | <b>26</b> |
| 10.1      | The effective mass tensor . . . . .                                                      | 26        |
| 10.2      | $\langle L \rangle$ , $\langle S \rangle$ and $\langle J \rangle$ on each atom . . . . . | 27        |
| 10.3      | The Fermi-surface nesting function . . . . .                                             | 28        |
| 10.4      | Magnetocrystalline anisotropy . . . . .                                                  | 29        |
| 10.5      | X-ray and magnetic structure factors . . . . .                                           | 31        |
| <b>11</b> | <b>Things with no <code>pw.x</code> counterpart</b>                                      | <b>32</b> |
| <b>12</b> | <b>Performance: dials, memory and GPUs</b>                                               | <b>33</b> |
| <b>13</b> | <b>Accuracy summary</b>                                                                  | <b>34</b> |
|           | <b>Where to go next</b>                                                                  | <b>35</b> |

# 1 Input and invocation

defumat reads **pw.x input files**. The same file that runs in Quantum ESPRESSO runs here, which is what makes every comparison in this document possible.

```
from defumat import Calculator

calc = Calculator.from_file("scf.in", pseudo_dir="pseudo", conv_thr=1.0e-8)
result = calc.get_scf()

print(result.total_energy, "Ry in", result.iterations, "iterations")
```

conv\_thr means what it means in a pw.x input: it is compared against QE's dr2, the Hartree energy of the density residual.

It need not be given at all. from\_file and from\_text **adopt the input's own &electrons namelist** — conv\_thr, mixing\_beta, mixing\_mode, mixing\_fixed\_ns and electron\_maxstep — so an input file that states how to converge itself converges the same way here as under pw.x. A keyword argument given to the constructor still wins over the file, which is how the snippet above overrides it. One variable of that namelist is read by pw.x and deliberately *not* adopted: diagonalization, because this package offers one eigensolver, so honouring diagonalization = 'cg' would turn a run that works into an error, and mapping it silently onto Davidson would be the kind of substitution refused everywhere else here. Name a solver at construction instead.

## One object, bound methods

A Calculator is a system together with its pseudopotentials, and every calculation in this document is a method on it. The pseudopotentials are read from the names the ATOMIC\_SPECIES card gave, resolved against pseudo\_dir — which from\_file defaults to the input file's own directory — so a script names an input and then names physics:

```
calc = Calculator.from_file("scf.in", pseudo_dir="pseudo")

bands = calc.get_bands(kpoints=path)      # runs the SCF first if none is cached
dos = calc.get_dos(grid=(8, 8, 8))
forces = calc.get_forces()
epsilon = calc.get_dielectric_tensor()
phonons = calc.get_phonons()

bands.plot()                             # and each result draws itself
```

Three properties are worth knowing before relying on it.

**The ground state is cached, and computed on demand.** A method that needs one and finds no cache runs an SCF, printing a line to stderr saying so; get\_scf() is the explicit path and is the same code. The cache is a *single slot* holding the result together with the options that produced it — calling get\_scf again with different options reruns and replaces it — because what it holds is the wavefunctions, and a cache keyed by option sets would quietly hold several sets of them. calc.scf\_result reads the slot without triggering anything, and a state that did not converge is refused rather than differentiated.

**Nothing mutates.** with\_positions, with\_cell and with\_spin return a *new* calculator with an

empty cache; the converged state crosses as a starting guess where that is sound, which for `with_spin` is the spin-regime continuation described under *Spin, magnetism and spin-orbit* and converges in one iteration rather than twenty-five. A cached result under a moved atom would otherwise answer for the geometry it was converged at.

**The refusals pass through untouched.** Every restriction in this document — PAW Born charges, a metal’s Raman tensor, a spiral with spin-orbit coupling — raises through the bound method exactly as through the function beneath it.

The functional entry points the rest of this document names — `run_scf`, `run_bands`, `dielectric_tensor` and the others — are unchanged and remain the way to drive a calculation whose state is being managed by hand:

```
from defumat.io.pwin import read_pw_input
from defumat.pseudo import read_upf
from defumat.scf.driver import Calculation, run_scf
from defumat.system import build_system

system = build_system(read_pw_input("scf.in"))
pseudos = tuple(read_upf(f"pseudo/{s.pseudo_file}") for s in system.structure.species)
result = run_scf(system, pseudos, conv_thr=1.0e-8)
```

What such a call has to carry with the density is the rest of the converged state — `becsum` for a PAW dataset, `ns` under a Hubbard  $U$ , `tau` under a meta-GGA. None of the three can be rebuilt from the density, and the entry points refuse without them rather than computing something else. The calculator passes them.

A command line covers the common workflows:

```
python3 -m defumat.cli inspect <qe-output> # summarise what the parser reads
python3 -m defumat.cli dos <input> # density of states
python3 -m defumat.cli pdos <input> # projected DOS
python3 -m defumat.cli relax <input> # structural relaxation
python3 -m defumat.cli stress <input> # stress tensor
python3 -m defumat.cli spiral <input> # spin-spiral scan
```

Standard `pw.x` variables — `ecutwfc`, `ecutrho`, `ibrav`, `nbnd`, `occupations`, `degauss` and the rest — mean what they mean there and are not re-documented here. A handful of knobs are worth knowing because they are specific to this code or to a feature below:

|                                                   |                                                           |
|---------------------------------------------------|-----------------------------------------------------------|
| <code>mbj_c</code>                                | fixes Tran-Blaha’s $c$ instead of taking the cell average |
| <code>london_s6</code> , <code>london_rcut</code> | Grimme D2’s scaling and real-space cutoff                 |
| <code>cell_dofree</code>                          | which cell degrees of freedom a vc-relax may move         |
| <code>starting_ns_eigenvalue</code>               | seeds the DFT+U occupation matrix                         |
| <code>spiral_q</code>                             | the spin-spiral wavevector, in lattice coordinates        |
| <code>LOCAL_MAGNETIC_FIELDS</code>                | a card with no <code>pw.x</code> counterpart              |

**Units are Rydberg atomic units throughout** (Ry, bohr), matching QE. The only layer that speaks anything else is `io/`: Hubbard  $U$  in the HUBBARD card is in eV and is converted at the boundary, as `pw.x` does it.

**Refuses.** `K_POINTS gamma` selects the half-sphere storage of the gamma-point trick, which is generated but not consumed. Such a run is substituted by an explicit  $k = 0$  on the full sphere — the same physics at twice the storage — and it says so.

## 2 Ground state

The Kohn-Sham problem, generalised because ultrasoft and PAW make the overlap non-trivial:

$$\begin{aligned}\hat{H} |\psi_{n\mathbf{k}}\rangle &= \epsilon_{n\mathbf{k}} \hat{S} |\psi_{n\mathbf{k}}\rangle, & \hat{S} &= 1 + \sum_{I,ij} q_{ij} |\beta_i^I\rangle \langle \beta_j^I| \\ \hat{H} &= -\nabla^2 + v_{\text{loc}}(\mathbf{r}) + v_{\text{H}}[n] + v_{xc}[n] & & + \sum_{I,ij} D_{ij}^I |\beta_i^I\rangle \langle \beta_j^I|\end{aligned}$$

with  $n(\mathbf{r}) = \sum_{n\mathbf{k}} w_{n\mathbf{k}} f_{n\mathbf{k}} (|\psi_{n\mathbf{k}}|^2 + \sum_{ij} Q_{ij}^I \langle \psi | \beta_i \rangle \langle \beta_j | \psi \rangle)$ , the augmentation term being what ultrasoft adds.  $D_{ij}^I$  is rebuilt from the potential every iteration, which is why the SCF is a fixed point in  $(n, D)$  and not in  $n$  alone. Convergence is measured as QE measures it — the Hartree energy of the residual:

$$dr^2 = \iint \frac{\delta n(\mathbf{r}) \delta n(\mathbf{r}')}{|\mathbf{r} - \mathbf{r}'|} d\mathbf{r} d\mathbf{r}' < \text{conv\_thr}$$

```
result = run_scf(system, pseudos, conv_thr=1.0e-10, mixing_mode="anderson")
print(result.total_energy, result.energy_terms) # QE's own term-by-term split
print(result.fermi_energy, result.homo, result.lumo)
```

**Metals** work with every smearing QE has (gaussian, marzari-vanderbilt, methfessel-paxton, fermi-dirac) and with the tetrahedron family — both linear and Bloechl-corrected — usable as an occupation scheme *inside* the SCF, not only for a density of states.

**Continuation across a change of spin regime.** Give `run_scf` a previous result as `starting_from` — `System.with_spin` underneath — and it promotes a converged state into a target's variables: non-magnetic into collinear, collinear into noncollinear, spin-orbit switched on. It reaches the *same* solution as a fresh run ( $\leq 4\text{e-}8$  Ry on six cases) in as little as one iteration. The magnetization is **seeded** from the target's `starting_magnetization` when the source has none, because nothing in the SCF breaks spin symmetry on its own and an unseeded promotion would converge back to the unpolarized solution and report success.

**Convergence.** `mixing_mode` selects `anderson` (default), `linear`, `broyden`, Kerker/Thomas-Fermi preconditioning (TF), or `local-TF`; `scf_solver` offers a Newton-Krylov alternative that is slower but reaches solutions the mixer cannot hold.

TF screens the residual with one Thomas-Fermi length for the whole cell. `local-TF` makes that length a function of  $\rho(\mathbf{r})$ , which is what a **slab** needs: the metal wants strong screening and the vacuum wants none, and a single compromise value over-screens one and under-screens the other — charge sloshing between the surfaces. The screened residual is the solution of

$$4\pi e^2 v(\mathbf{G}) + |\mathbf{G}|^2 (\alpha v)(\mathbf{G}) = |\mathbf{G}|^2 (\alpha \delta n)(\mathbf{G}), \quad \alpha(\mathbf{r}) = 3 \left( \frac{2\pi}{3} \right)^{5/3} r_s(\mathbf{r}),$$

with  $(\alpha f)$  multiplied in *real* space, so the operator is not diagonal in  $\mathbf{G}$  and is inverted iteratively at each step. On a cobalt film it converges at `mixing_beta` = 0.7 where plain Anderson at the same  $\beta$  diverges and Kerker at 0.3 is still far out after 250 iterations. `max_iterations` defaults to 100 — **a hard SCF may need more**, and hitting the cap is reported as not converged rather than as an answer.

### 3 Bands, densities of states, projections

A band structure is one diagonalisation per  $\mathbf{k}$  on a fixed density — `run_bands`, or `run_nscf` for a non-self-consistent run on a grid rather than a path. A density of states is

$$g(E) = \sum_{n\mathbf{k}} w_{\mathbf{k}} \delta(E - \epsilon_{n\mathbf{k}})$$

with  $\delta$  replaced by a smearing function or evaluated by tetrahedron interpolation. The projected DOS weights each term by  $|\langle \phi_\mu | \hat{S} | \psi_{n\mathbf{k}} \rangle|^2$ , and the spilling parameter is what is *not* captured,  $\sum_{n\mathbf{k}} w_{\mathbf{k}} f_{n\mathbf{k}} (1 - \sum_\mu |\langle \phi_\mu | \hat{S} | \psi \rangle|^2)$ .

```
from defumat.workflows.bands import run_bands
from defumat.workflows.dos import run_dos
from defumat.workflows.pdos import run_pdos

bands = run_bands(system, pseudos, scf.density, kpoints=path)  # a k-path

dos, _ = run_dos(system, pseudos, scf.density, grid=(12, 12, 12),
                 scheme="tetrahedra_opt")  # -> (DensityOfStates, NSCFResult)

pdos, _ = run_pdos(system, pseudos, scf)  # -> (ProjectedDOS, NSCFResult)
print(pdos.channels[0])  # resolved by atom, l and m
print(pdos.charges.charges, pdos.charges.spilling)
```

The DOS schemes live behind a name registry (`DOS_SCHEMES`), and the projected DOS feeds *the same* registry as a per-band weight, so a PDOS and a DOS are the same integration. Projections are  $\langle \phi | \hat{S} | \psi \rangle$  on Löwdin-orthogonalised pseudo-atomic orbitals. Validated against a purpose-built `projwfc.x` on seven cases: **6.9e-4** on a projection, **4.7e-5** on a Löwdin charge, and the spilling parameter to all four decimals it prints.

### 4 Pseudopotentials and functionals

**Kinds:** norm-conserving, **ultrasoft** and **PAW** — the two-grid split, the augmentation charge, the overlap operator, self-consistent  $D_{ij}$ , and PAW’s one-centre terms including its radial Poisson solve and spherical quadrature. UPF v2 is read.

**Functionals:** LDA (PZ), **PBE**, **revPBE**, **PBEsol**, on the plane-wave grid and on the PAW spheres. The functional comes from the pseudopotential headers unless `input_dft` overrides it. Only the *energy* is written down;  $v_{xc}$ , and a GGA’s  $v_1$  and  $v_2$ , come from `jax.grad` of it.

**Meta-GGA, potential-only branch:** `input_dft = 'tb09'` (Tran-Blaha) and `'bj06'` (Becke-Johnson). These invert the rule above — they *are* potentials, there is no energy — so the consequences are enforced rather than documented: `run_scf` warns that its total is not the value of any functional it minimised, and forces, stress, phonons and response all refuse. Silicon’s gap goes from LDA’s 0.49 eV to **1.13 eV** against an experimental 1.17.

**Van der Waals:** Grimme **D2** (`vdw_corr = 'grimme-d2'`). Bilayer graphene matches `pw.x` to **3.1e-9 Ry** and binds at 3.23 Å where PBE alone has no minimum.

**Refuses.** UPF v1; an unimplemented functional (rather than silently substituting LDA); energy-carrying meta-GGAs (TPSS, SCAN, M06L), whose potential has a  $\partial E / \partial \tau$

piece acting on the wavefunction that is not written; plain *ultrasoft* under a meta-GGA, which has no partial waves to reconstruct  $\tau$  from inside the sphere where PAW does; EXX; and **D3, Tkatchenko-Scheffler, MBD and XDM** — where `pw.x` warns and silently runs with no correction at all.

## 5 Spin, magnetism and spin-orbit

| regime       | how to ask                       | note                                                         |
|--------------|----------------------------------|--------------------------------------------------------------|
| unpolarised  | default                          |                                                              |
| collinear    | <code>nspin = 2</code>           | one Fermi level, or two under <code>tot_magnetization</code> |
| noncollinear | <code>noncolin = .true.</code>   | magnetization is a vector                                    |
| spin-orbit   | <code>+ lspinorb = .true.</code> | $j$ -resolved projectors; NC, US and PAW                     |

**Three spin numbers are kept apart**, and `System` exposes all three: `nspin` says which regime is in force, `npol` is the number of spinor components of a *wavefunction*, and `nspin_mag` the number of components of a *density*. They coincide at 1 and 2 and come apart at 4, where `npol` = 2 but `nspin_mag` is **one** for a nonmagnetic spin-orbit run — which is why such a run costs about what a doubled unpolarized one costs.

**Fixed occupations under `nspin = 2`** fill each channel to its own count rather than searching for a level: `tot_magnetization` gives `nelup` and `neldw`, and `iweights_only` fills `NINT( $n_{\uparrow}$ )` bands in one channel and `NINT( $n_{\downarrow}$ )` in the other. That is the *only* shape the combination has — `pw.x` refuses fixed-occupation LSDA without a `tot_magnetization`, and requires an integer one — and this package makes the same two refusals at input, before anything is allocated. Each channel’s highest occupied level is reported as `fermi_energy_up` / `fermi_energy_down` beside the overall `homo` and `lumo`, which is what `iweights` returns. Validated on an oxygen atom at `tot_magnetization` = 2, whose four-up / two-down filling is the Hund’s-rule ground state: **4.9e-9 Ry** against `pw.x`.

**Magnetic fields and constraints:** a uniform field over the cell, a field inside one atom’s sphere (through a `LOCAL_MAGNETIC_FIELDS` card that `pw.x` has no counterpart for), Elk’s `reducebf`, and all four of QE’s `constrained_magnetization` schemes. **The field’s energy is not part of the reported total** — QE prints `etcon` and never adds it, and Elk excludes its external field’s energy by the same convention.

**DFT+U:** both rotationally-invariant functionals, on atomic, ortho-atomic and norm-atomic projectors, read from the `HUBBARD` card. Forces come from differentiating through projectors that move with the atoms.

*The card selects which functional runs*, exactly as `pw.x` does it — the `namelist`’s `lda_plus_u_kind` is obsolete there and warned about. A card carrying only  $U$ ,  $J_0$ , `ALPHA` and `BETA` is Dudarev’s simplified functional; any  $J$ ,  $B$ ,  $E_2$  or  $E_3$  selects Liechtenstein’s full one, whose interaction is the whole Coulomb matrix

$$\text{vee}(m_1 m_2 m_3 m_4) = \sum_k F^k \frac{4\pi}{2k+1} \sum_{q=-k}^k a_{kq}(m_1 m_3) a_{kq}(m_2 m_4)$$

rather than one number, with the Slater integrals  $F^k$  fixed by  $U$  and  $J$ . Against `pw.x` on antiferromagnetic FeO: **4.2e-9 Ry** at  $J$  = 1 eV. *Forces work for both*, and for the full functional `pw.x` does not have them at all (`force_hub.f90` stops on “forces in the DFT+U+J scheme are not implemented”); a central difference of the SCF energy agrees to 1.3e-4 Ry/bohr. **Seed that difference from the undisplaced run** — with a full interaction

matrix the DFT+U landscape has more than one solution, and cold-started displacements land on a different one, which does not look like an error, it looks like a force.

```
from defumat import Calculator
scf = Calculator.from_file("feo.in").get_scf()
print(scf.energy_terms["hubbard"])
```

```
HUBBARD {atomic}
U Fe1-3d 4.3
J Fe1-3d 1.0      # selects the full functional
```

**Around-mean-field double counting** (`hubbard_double_counting = 'amf'`) is the alternative to the fully-localised limit that both QE's functionals assume, and `pw.x` has neither the option nor the code. FLL treats the occupations as integers, which is right for a localised magnetic insulator and wrong for a metal; AMF subtracts the shell's mean occupation before the interaction, so a uniformly filled shell is corrected by *exactly* nothing — 0 to  $10^{-14}$  at any filling, where FLL is not. On FeO its Hubbard energy is  $-0.005$  Ry against FLL's  $+0.232$ . It needs the full functional, there being no interaction matrix in the simplified one to apply the shift to.

**Slater integrals from the orbital** (`hubbard_slater = 'yukawa'`) computes the interaction instead of parameterising it:  $F^k$  is the double radial integral of the manifold's own all-electron partial wave against a screened Coulomb kernel,

$$F^k = \iint dr dr' [r\phi(r)]^2 [r'\phi(r')]^2 (2k+1) \lambda i_k(\lambda r_<) \tilde{k}_k(\lambda r_>),$$

which becomes the bare  $r_<^k/r_>^{k+1}$  as  $\lambda \rightarrow 0$ . Give  $\lambda$  on the card (`LAMBDA Ni-3d 2.1`) or give a  $U$  and let the screening length be solved for — either way  $F^0$ ,  $F^2$ ,  $F^4$  and  $J$  all come out of the orbital. Nickel's unscreened 3d gives  $F^4/F^2 = 0.626$  against the 0.625 that QE and Elk both hardcode, and  $J = 1.28$  eV.

```
&system
  hubbard_slater = 'yukawa'
/
HUBBARD {ortho-atomic}
U Ni-3d 5.0      # lambda is solved for; F2, F4 and J follow
```

**On a spinor** (`noncolin`, with or without `lspinorb`) the occupation matrix is one  $2 \times 2$  matrix in spin space rather than two real matrices, measured on projector columns that are themselves spinors, and its off-diagonal blocks are what let the shell's moment point anywhere. Both functionals work there. Against `pw.x`: **7.7e-9 Ry** on relativistic BN and **1.4e-9 Ry** on fully-relativistic iron carrying  $U = 2.20$ ,  $J = 1.75$  eV with its moment turned into the  $xy$  plane, with  $\text{Tr}[n^{\sigma\sigma}]$  matching every digit `write_ns` prints.

**Tensor moments** put that matrix in a basis where each element is a multipole rather than a number: an orthonormal, complete set  $\Gamma_t^{kpr}$  with  $k$  the rank in the orbital index,  $p$  the rank in spin and  $r$  their coupling, so that

$$n = \sum_{kp rt} w_t^{kpr} \Gamma_t^{kpr},$$

with real  $w$ . The charge is  $w^{000}$ , the spin moment is  $w^{011}$ , and  $\mathbf{L} \cdot \mathbf{S}$  is  $w_0^{110}$  alone — to  $10^{-11}$  over all one hundred moments of a  $d$  shell.

**Fixing one** (a `TENSOR_MOMENTS` card) converges the field to a state carrying it, which selects an orbital or multipolar ordering the field would not find on its own. The penalty is written



down and its potential is `jax.grad` of it, and, like a constrained moment's, its energy is not part of the reported total — so the number that comes back is the unconstrained functional at the constrained state.

```
&system
  tensor_moment_penalty = 20.0
/
TENSOR_MOMENTS
N-2p 1 1 0 0 -0.40      # hold L.S at -0.40
```

Being quadratic, the penalty settles where its gradient balances the material's own restoring force, so the moment *approaches* its target as the penalty stiffens rather than reaching it: on nitrogen's  $2p$ , whose free  $\mathbf{L} \cdot \mathbf{S}$  is  $-0.0006$ , holding it at  $-0.40$  gives  $-0.109$ ,  $-0.263$ ,  $-0.355$  and  $-0.388$  at  $\lambda = 1, 5, 20, 80$ , with the energy rising monotonically by 2.6, 15.3, 27.4 and 32.6 mRy.

**Refuses.** A spinor occupation matrix needs `nosym`: its group average takes the  $SU(2)$  representation of each operation beside the rotation of the  $m$  indices, and averaging the  $m$  indices alone leaves the moment pointing somewhere else. A fixed tensor moment needs a spinor: for a collinear run the spin-rank moments keep only their  $z$  component and the constraint degenerates into the fixed spin moment `constrained_magnetization` already is. `hubbard_slater = 'yukawa'` needs a **PAW** dataset, and this is a measurement rather than a preference: only PAW carries the all-electron partial wave. A norm-conserving  $\chi$  is normalised and gives a pseudo-orbital  $F^0$  far too small (silicon's  $3p$ : 8.8 eV), and an *ultrasoft*  $\chi$  is not even normalised — iron's  $3d$  has  $\int (r\chi)^2 = 0.41$ , the rest of the norm living in the augmentation charge, and its  $F^0$  comes out at 2.1 eV where the answer is above twenty. The `norm` that comes back with each manifold is what says whether it is bound inside the cutoff, and the cutoff defaults to the dataset's own augmentation radius because a partial wave *diverges* outside it. Elk's `inpdfu = 2` and `3` (Slater and Racah *input*) have no card syntax: they are a change of coordinates on the three numbers  $(U, J, B)$  already spans, and QE's card has spent the names  $E_2$  and  $E_3$  already. The FLL/AMF interpolation (Elk's `dfu = 3`) is not here because it is not in Elk either any more — documented in the 11.0.2 manual, removed from the 11.0.2 source.

**Refuses.** The intersite  $V$ , background channels, the orbital-resolved variant, the `wf/pseudo` projector sets, and noncollinear `ns_nc`; and PAW + GGA + a noncollinear magnetization together.

A **magnetic field in a noncollinear run with no starting moment** is refused rather than run. `domag` comes from `starting_magnetization` and from nothing else (`setup.f90`), so such a run has `nspin_mag = 1` and a density with no magnetization channels for the field to act on. Set `starting_magnetization` for one species — any nonzero value will do, the field decides where the moment goes. Worth knowing: `pw.x` *accepts* this input and silently drops the field, because `add_bfield` writes it into the potential's magnetization channels and `vloc_psi_nc` applies those only IF (`domag`).

The **residual solver** cannot take a fixed occupation whose filled/empty boundary cuts a **degenerate multiplet** — which is exactly what an atom's Hund's-rule filling does. Which member of the multiplet the eigensolver returns is arbitrary, so the density map is not a function of the density and a Newton step has nothing to converge on; it is diagnosed by name rather than left as a NaN. The mixer damps its way past it, and a *gapped* fixed-occupation LSDA cell agrees between the two solvers to  $5e-12$  Ry.

## 6 Forces, stress and relaxation

$$\mathbf{F}_I = -\frac{\partial E}{\partial \tau_I} \Big|_{\psi \text{ fixed}}, \quad \sigma_{ij} = -\frac{1}{\Omega} \frac{\partial E}{\partial \varepsilon_{ij}} \Big|_{\psi \text{ fixed}}$$

Both are one gradient. With ultrasoft the orthonormality constraint  $\langle \psi_n | \hat{S} | \psi_m \rangle = \delta_{nm}$  is carried explicitly, so its multiplier term — QE’s Pulay contribution — falls out of the same derivative rather than being added afterwards. A variable-cell relaxation minimises the **enthalpy**, which is why a relaxed crystal carries the applied pressure rather than having zero stress:

$$H = E + P \Omega, \quad \frac{\partial H}{\partial h} = \Omega (P \mathbf{1} - \sigma) h^{-T}$$

```
from defumat.forces import compute_forces
from defumat.stress import compute_stress
from defumat.workflows.relax import run_relax
from defumat.workflows.vc_relax import run_vc_relax

forces = compute_forces(calculation, result) # method=... for QE's own terms
print(forces.forces) # (nat, 3) in Ry/bohr

stress = compute_stress(calculation, result, terms=True)
print(stress.pressure) # kbar

relaxed = run_relax(system, pseudos, forc_conv_thr=1.0e-4)
cell = run_vc_relax(system, pseudos, press=500.0) # kbar
print(cell.pulay_error) # the frozen-basis error, reported not hidden
```

QE’s `force_lc`, `force_cc`, `force_ew`, `force_us`, `addusforce` and `force_corr` are transcribed too, behind the same registry — **as a cross-check, not as the implementation**. The two share no machinery, which is what makes disagreement informative; it is how the augmentation force’s sign and a missing gradient correction were found.

For a noncollinear or spin-orbit run the same expression is evaluated on two-component spinors: the nonlocal term takes the pseudopotential’s bare `dvan_so` (a complex  $2 \times 2$  matrix in spin space) and the orthonormality constraint takes `qq_so`, and the state is one coefficient vector of length  $2n_{\text{pw}}$  rather than two. Nothing else in the functional changes.

```
from defumat import Calculator

calculator = Calculator.from_file("pt2-soc-force.in") # noncolin, lspinorb
forces = calculator.get_forces() # (nat, 3) Ry/bohr
stress = calculator.get_stress() # (3, 3) Ry/bohr^3
print(forces.max_force, stress.pressure_kbar)
```

Validated against `pw.x` with `tpnfor/tstress` on a four-atom noncollinear hydrogen chain (norm-conserving) to  $8.9 \times 10^{-7}$  Ry/bohr, doubled fcc platinum with spin-orbit coupling to  $7.5 \times 10^{-6}$  (ultrasoft) and  $7.3 \times 10^{-7}$  (PAW), and the stress on those three plus QE’s own three `pw_spinorbit` cases to  $\leq 1.2 \times 10^{-6}$  Ry/bohr<sup>3</sup>; and, without any Fortran at all, against a central finite difference of the frozen energy to  $6.2 \times 10^{-9}$  Ry/bohr.

A noncollinear or spin-orbit run has forces, a stress and a relaxation, but only through the *autodiff* method: `method='analytic'` is refused by name, because `force_us/stres_knl`

are transcriptions of QE's hand-derived expressions and have no spinor form. The force on an atom of a **spin spiral** is refused separately — its two spinor components live on different plane-wave spheres, so the nonlocal term needs the projectors of both;  $dE/dq$  is what a spiral has instead. **Elastic constants and electrostriction** stay refused for `noncolin`: they differentiate the same functional a third time and their first-order wavefunctions come from a Sternheimer solve that has no spinor form.

## 7 Response: dielectric, phonons, Raman

The **velocity operator** comes first and everything else uses it: it is one jvp of  $\hat{H}(\mathbf{k})$  at a frozen sphere, because  $\partial\hat{H}/\partial k_\alpha = i[\hat{H}, r_\alpha]$  in the periodic gauge. `band_velocities` exposes it — `Calculator.get_band_velocities` is the short way — and because the overlap carries a velocity too, a band velocity is  $\langle\psi|\partial_k\hat{H} - \epsilon\partial_k\hat{S}|\psi\rangle$ .

```
calc = Calculator.from_file("scf.in", pseudo_dir="pseudo")

velocities = calc.get_band_velocities()           # .velocities is (nk, nbnd, 3)
elsewhere = calc.get_band_velocities(kpoints=path)
```

Every first-order quantity then comes from the Sternheimer equation — the response of an occupied orbital to a perturbation, projected outside the occupied manifold so that no empty states and no energy denominators appear:

$$(\hat{H} - \epsilon_n\hat{S} + \alpha\hat{Q})|\Delta\psi_n\rangle = -\hat{P}_c^+ \Delta V|\psi_n\rangle$$

It is solved self-consistently, because  $\Delta V$  contains the response of the potential to the response of the density,

$$\Delta V_{\text{scf}} = \Delta V_{\text{ext}} + \int K(\mathbf{r}, \mathbf{r}') \Delta n(\mathbf{r}') d\mathbf{r}'$$

and the kernel  $K$  is one jvp of `v_of_rho` rather than a second implementation. From its solutions:

$$\epsilon_{\alpha\beta}^\infty = \delta_{\alpha\beta} + \frac{8\pi}{\Omega} \sum_{\mathbf{k}} w_{\mathbf{k}} \langle \Delta_\alpha \psi_n | \partial_\beta \psi_n \rangle, \quad Z_{I,\alpha\beta}^* = \frac{\partial F_{I\alpha}}{\partial \mathcal{E}_\beta}$$

$$C_{I\alpha,J\beta} = \frac{\partial^2 E}{\partial \tau_{I\alpha} \partial \tau_{J\beta}}, \quad \omega^2 \mathbf{e} = \frac{1}{\sqrt{M_I M_J}} C \mathbf{e}, \quad \frac{\partial \epsilon_{ij}}{\partial \tau_{I\gamma}} \text{ (Raman)}$$

The Raman tensor is the second-order energy differentiated once more along a displacement, at *frozen first-order* wavefunctions — the  $2n+1$  theorem, which is why a third derivative costs one more jvp and not a new solve.

```
from defumat.response import (dielectric_tensor, dynamical_matrix,
                              raman_tensors, vibrational_spectrum)

eps = dielectric_tensor(calculation, scf.wavefunctions, scf.eigenvalues,
                        scf.density)
print(eps.epsilon)           # 3x3, cartesian

ph = dynamical_matrix(calculation, scf.wavefunctions, scf.eigenvalues,
```

```

scf.density, scf.becsum)
print(ph.frequencies)          # cm-1

raman = raman_tensors(calculation, scf, born_charges=True, keep_internals=True)
spectrum = vibrational_spectrum(calculation, scf) # per-mode Raman and IR
print(spectrum.raman_activity, spectrum.depolarisation)

```

born\_effective\_charges gives  $Z^*$  on its own, and mode\_activities is the bare assembly — a transcription of dynmat.x's RamanIR, kept a pure function of arrays so it is testable without a self-consistent run.

Ultrasoft and PAW work for the dielectric constant ( $\leq 1.2\text{e-}4$  against ph.x on four cases) *and for the dynamical matrix* — silicon's optical mode comes out at **513.295** and **513.378**  $\text{cm}^{-1}$  against ph.x's 513.275 and 513.404, which is tighter than the norm-conserving case's 0.05. Four things make the difference and all four switch themselves off when  $S$  is the identity: the source term is  $(dH/du - \epsilon dS/du)|\psi\rangle$ ; the first-order state has an occupied block the Sternheimer solve does not produce; the mixed state changes at *frozen* states, because the augmentation charge and the projectors travel with their atom; and the orthonormality multipliers are variables of the functional that move — as a **matrix**, since a diagonal one is not invariant under the occupied-manifold rotation the state tangent is free in. Metals work for the response and for a norm-conserving dynamical matrix; and a symmetry-reduced wedge works as of the rank- $n$  symmetriser.

**The Raman tensor is ultrasoft and PAW as well**, at **1.2e-4** on both against a central difference of  $\epsilon$  over re-converged displaced cells, where the norm-conserving control on the same script is  $6.8\text{e-}4$ . It took **two tangents that are only right together**, and either alone is worse than neither: the state tangent is  $P_c d\psi + \psi^{\text{ort}}$ , the occupied block again; and  $|b\rangle = P_c \mathbf{r}|\psi\rangle$  is *not* the solution of its own linear equation once adddvpsi\_us has applied  $S$  to it and added the augmentation dipole, so  $db$  is the tangent of a *composition* rather than of a solve. What located them is that  $\partial\epsilon/\partial\tau$  is a sum of five partial derivatives and each one can be finite-differenced on its own against a re-converged run: three agreed to  $7\text{e-}4$  and two did not.

**A trap worth knowing:** a response on a reduced  $k$ -set is a *polar vector field* and must be symmetrised as one. Running the whole grid instead is sound only if the grid is closed under the point group — a **shifted** Monkhorst-Pack grid is not.

**Refuses.** Born charges for PAW (int3\_paw against becsumort is missing, and without it the expression is wrong in sign as well as size);  $\chi^{(2)}$  and the electro-optic tensor (the second-order source term has nothing to build the  $\langle u_i | r_k | u_j \rangle$  piece from — and it is **42%** of the answer, measured); phonons at  $q \neq 0$ ; the dynamical matrix of an ultrasoft or PAW *metal* (the wg/wk weight split was derived for a response whose becsum dependence is entirely inside dps\_i, and an insulator cannot say which weight the three further tangents belong with); any response under a potential-only meta-GGA; and a shifted grid with a reduced  $k$ -set.

## LO-TO splitting and the static dielectric constant

At  $\Gamma$  a polar crystal's dynamical matrix is not analytic. An optical mode that moves the ions against each other builds a macroscopic electric field, and what the field costs depends on the *direction* the zone centre is approached from, so the matrix acquires a rank-one

term:

$$D_{I\alpha,J\beta} \rightarrow D_{I\alpha,J\beta} + \frac{4\pi e^2}{\Omega} \frac{(\mathbf{q} \cdot \mathbf{Z}_I^*)_\alpha (\mathbf{q} \cdot \mathbf{Z}_J^*)_\beta}{\mathbf{q} \cdot \epsilon^\infty \cdot \mathbf{q}}$$

It raises the longitudinal mode and leaves the transverse ones alone — the LO-TO splitting — and it vanishes identically where  $\mathbf{Z}^*$  does, so a non-polar crystal is untouched. Contracting the same two ingredients the other way gives the ionic screening of a *static* field, one oscillator per polar mode:

$$\epsilon_{ij}^0 = \epsilon_{ij}^\infty + \frac{4\pi e^2}{\Omega} \sum_\nu \frac{p_i^\nu p_j^\nu}{\omega_\nu^2}, \quad p_i^\nu = \sum_{I\gamma} Z_{I,i\gamma}^* z_{I\gamma}^\nu$$

where  $p$  is the same mode dipole the infrared activity is the square of. So  $\epsilon^0$  and  $\epsilon^\infty$  are two ends of one contraction, and together with the frequencies they satisfy the Lyddane-Sachs-Teller relation  $\epsilon^0/\epsilon^\infty = (\omega_{\text{LO}}/\omega_{\text{TO}})^2$ .

```
spectrum = calc.get_vibrational_spectrum(neutralize=True)
print(spectrum.static_permittivity)           # eps_0, 3x3
print(spectrum.ionic_permittivity)           # the polar modes' share of it

lo = calc.get_vibrational_spectrum(loto_direction=(1, 0, 0), neutralize=True)
print(lo.frequencies)                         # the LO modes along x
```

`nonanal` adds the term to a dynamical matrix on its own and `loto_modes` diagonalises the result; `polar_mode_permittivity` is the oscillator sum, and `neutral_born_charges` the sum-rule correction below. All four are pure functions of arrays, so they can be checked without a self-consistent run behind them.

Both halves agree with `dynmat.x` — which applies the same term and the same oscillator sum to the tensors this code writes to its `fildyn` — to every digit it prints. **The Lyddane-Sachs-Teller relation is the check neither half can pass alone**, since one side is a ratio of two diagonalisations and the other a sum of dipoles over  $\omega^2$ ; on AIs it holds to **5.0e-11**.

**It needs a charge-neutral  $\mathbf{Z}^*$  and that is worth knowing before reading an LO frequency.** The sum rule  $\sum_I \mathbf{Z}_I^* = 0$  says a rigid translation builds no field; a computed  $\mathbf{Z}^*$  misses it by the basis-set error, which *charges the crystal* — AIs at `ecutwfc = 10` misses it by  $-1.257$ , and the non-analytic term then lifts a longitudinal *acoustic* mode from  $1.8$  to  $33.8 \text{ cm}^{-1}$  and puts the LO frequency  $7.7 \text{ cm}^{-1}$  low. `neutralize=True` shares the violation out over the atoms, which is what `dynmat.x` does under any `asr` but `'no'`. Comparing against another code does not see this: it is handed the same charges and reproduces the same wrong number. LST does.

**Refuses.** Everything the Raman tensors and the Born charges refuse, since it is built from both: PAW  $\mathbf{Z}^*$ , metals, and the rest of the list above. The splitting is a  $\Gamma$  quantity by construction — it is what  $q \rightarrow 0$  leaves behind — so a dispersion through it needs phonons at  $q \neq 0$ , which are not implemented.

## Elastic constants, electrostriction and the elasto-optic tensor

The same machinery in the *strain* coordinate rather than the atomic one. **The strain response itself works on all three pseudopotential kinds** (`strain_response`): the augmentation charge  $Q_{ij}(\mathbf{r})$  is a function of the *cell*, so a homogeneous strain deforms the table it is tabulated on where a displacement only translates it, and on top of that come

the four terms the dynamical matrix needed. Against a central difference of *re-converged* strained runs — which needs no `ph.x`, since `ph.x` has no strain perturbation — it agrees to **4.6e-4** (ultrasoft) and **4.7e-4** (PAW) against a norm-conserving 1.9e-4 on the same cell. The *derivatives* below are norm-conserving still — but the variational second-order energy they all stand on is not: it reproduces `dielec.f90`’s dielectric constant to **3.4e-10** (ultrasoft) and **6.9e-11** (PAW) against a norm-conserving 8.4e-10, which took four terms that each vanish when  $S$  is the identity. One of them is worth stating on its own, because it is a distinction rather than a term: a *state* is projected with  $1 - \sum |\psi\rangle\langle\psi|S$  and a *right-hand side* with  $1 - \sum S|\psi\rangle\langle\psi|$ , and collapsing the two — which is exact at  $S = 1$  — costs a factor of fourteen. The third derivative in this *strain* coordinate is what stays norm-conserving, and how far off it is has been measured rather than left unknown: both of the tangents that closed the Raman tensor transfer, and take it from **4.6e-2** to **1.3e-2** against a central difference over re-converged strained cells, where the norm-conserving control on the same script is **2.3e-4**. What is left is entirely the *db* term — the same  $-1.72$  of 112 on ultrasoft and on PAW, which is what says it is structural. One candidate for it is excluded by measurement: writing the commutator that sources *db* with the multiplier matrix rather than the frozen scalar eigenvalue closes this coordinate (1.7e-4) and breaks the *displacement* one, taking the Raman tensor from 1.2e-4 to 1.14e-3. `strain_response` solves the first-order problem for a strain — Abinit’s metric-tensor formulation, obtained for nothing here because the cell was already written in reduced coordinates — and two things follow from it:

$$C_{ijkl} = \frac{\partial \sigma_{ij}}{\partial \varepsilon_{kl}}, \quad \frac{\partial \chi_{ij}}{\partial \varepsilon_{kl}} \quad (\text{elasto-optic})$$

The second is a **third** derivative of the energy, and it comes from one `jvp` of the second-order energy at frozen first-order wavefunctions — the  $2n+1$  theorem again. Through the thermodynamic identity of Tanner, Bousquet and Janolin it gives the four electrostriction tensors  $m$ ,  $q$ ,  $M$  and  $Q$ .

```
from defumat.response import elastic_constants, electrostriction

es = electrostriction(calculation, scf)      # solves the strain response itself
print(es.m, es.q, es.M, es.Q)              # the four electrostriction tensors
print(es.photoelastic)                     # the elasto-optic tensor
print(es.elastic.voigt)                    # C_ij in Voigt notation, GPa

# or the elastic constants alone, from a strain response you already have
C = elastic_constants(calculation, scf.wavefunctions, scf.eigenvalues,
                      scf.density, es.strain)
print(C.tensor.shape, C.voigt)
```

Converged silicon gives  $C_{11} = 198.5$  GPa and  $C_{12} = 68.9$ , against a measured 165.7 and 63.9; the three independent components of  $\partial\epsilon/\partial\varepsilon$  match a central difference over re-converged strained cells to **2e-4**, which is that difference’s own floor, and the rank-4 tensor is cubic to 3e-14 with nothing imposing it.

**Refuses.** Both are norm-conserving, `nspin = 1`, insulators and **clamped-ion**. A diverged first-order solution is refused rather than consumed in silence (`require_converged_responses`) — at QE’s `alpha_mix = 0.7` the strain response of a *slab* diverges, and consuming it gave a  $C_{ijkl}$  that was not even symmetric under  $C_{ijkl} = C_{klij}$ . The elastic constants additionally need the whole  $k$ -grid rather than a wedge, because their functional builds its own density and symmetrises it as a scalar inside the chain rule.



## 7.1 The piezoelectric tensor

The polarization a strain induces is the stress an electric field induces — one mixed second derivative read either way round:

$$e_{k,ij} = \frac{\partial P_k}{\partial \varepsilon_{ij}} = \frac{\partial \sigma_{ij}}{\partial \mathcal{E}_k} = -\frac{1}{\Omega} \frac{\partial^2 E}{\partial \varepsilon_{ij} \partial \mathcal{E}_k}$$

It is the Born charge's construction with one coordinate changed.  $Z^*$  is a jvp of the *force* along the field's response; the force is `jax.grad` of the frozen-state energy in the atomic positions and the stress is `jax.grad` of the same functional in a strain, so this is a jvp of the *stress* along the same response — three of them, one per field direction, on top of a dielectric constant that was going to be solved anyway. The orthonormality constraint contributes nothing here, because  $\langle \psi | \hat{S} | \psi \rangle$  is a sum over a sphere of integers and carries no cell.

```
calc = Calculator.from_file("alas.in", pseudo_dir="pseudo")

piezo = calc.get_piezoelectric_tensor()
print(piezo.voigt)           # (3, 6) in C/m^2, columns xx yy zz yz xz xy
print(piezo.e14)            # -0.7638: zincblende's one component
print(piezo.e.shape)        # (3, 3, 3) in e/bohr^2, e[k, i, j]
print(piezo.dielectric.epsilon) # the field response is shared, not re-solved
```

`pw.x` computes no piezoelectric tensor — the word occurs once in the whole distribution, in a citation in a comment — and Elk's task 380 reaches it by a finite difference of the Berry-phase polarization over one full ground state per strain tensor. So the validation is internal, and there are four independent statements. Silicon is centrosymmetric and its whole tensor vanishes (2.4e-5 C/m<sup>2</sup> against AlAs's 0.764 from the same code, with nothing imposing it on a `nosym` run); AlAs is  $\bar{4}3m$  and only  $e_{14} = e_{25} = e_{36}$  survive, to 1.7e-14; an eight-point wedge reproduces the sixty-four-point closed grid to **4.5e-9**; and the same mixed derivative contracted the other two ways — `zstar_eu.f90`'s expression with a strain label, and the strain response against the field's bare perturbation — agree to **6.2e-15** and **1.3e-7**. What anchors the sign and the normalisation is that the *same* assembly run in the position coordinate is the Born charge, and that is `ph.x`'s number.

**Refuses. Clamped-ion** — the atoms are carried by the strain and do not relax, which is also what Elk's task computes; the internal-strain term that makes the constant comparable with experiment is not here, and for zincblende it very nearly cancels this one. And **polar crystals by name**: what the derivative above gives is the *improper* tensor, and the proper piezoelectric response differs from it by  $\delta_{ki}P_j - \delta_{ij}P_k$ , which needs the spontaneous polarization itself. Those terms vanish identically whenever the two Cartesian labels they pair are different — so  $e_{14}$  of a zincblende crystal never carries the ambiguity — and they vanish for every component of a class that admits no invariant vector, which is what is checked. **Ultrasoft and PAW**, and that one is a gap rather than a missing term: nothing in the assembly is norm-conserving and the displacement leg of it is the Born charge, which is validated on all three kinds, but every ultrasoft and PAW crystal here is centrosymmetric and agrees with zero whatever is wrong. Beyond that: insulators, `nspin = 1`, and everything the Sternheimer regime and the differentiable cell refuse (a metal, a spin spiral, a magnetic field, a shifted grid run with `nosym`).

## 7.2 Spin-polarized response

The Sternheimer equation is solved per spin channel,

$$(H_\sigma - \varepsilon_{n\sigma}S + \alpha_{pv}Q_\sigma) |\Delta\psi_{n\sigma}\rangle = -P_{c,\sigma}^+ \Delta V_\sigma |\psi_{n\sigma}\rangle,$$

with the occupied manifold  $P_{c,\sigma}$  built from *that channel's* filling:  $N_{\uparrow}^{\text{occ}} = \text{NINT}(n_{\uparrow})$  and  $N_{\downarrow}^{\text{occ}} = \text{NINT}(n_{\downarrow})$ , which for a magnetic insulator are different numbers. Quantum ESPRESSO never meets this because LSDA doubles nks there and each spin channel is a separate  $k$ -point.

```
# chi_0 for a spin-polarized run -- the part that is validated unconditionally.
from defumat.response.sternheimer import make_sternheimer

solver = make_sternheimer(calculation, result, spin_polarized=True)
drho = solver.chi0(dv)          # dv is (2, n1, n2, n3) on the dense grid
```

Validated against a central difference of the density under a spin-dependent probe potential —  $1.8 \times 10^{-6}$  on an antiferromagnetic hydrogen chain (a smeared metal) and  $1.1 \times 10^{-6}$  on triplet O<sub>2</sub> (the sliced branch, ultrasoft) — and against the identity that a cell with no magnetization gives the same  $\varepsilon_\infty$  at `nspin = 1` and `nspin = 2`:  $6.2 \times 10^{-14}$  on 13.806646105. The probe must differ between the channels, because  $\chi_0$  is block diagonal in spin and one equal in both would not tell the blocks apart.

**Refuses.** `tot_magnetization` together with a smearing: the two channels have separate Fermi levels and `Smearing` carries one scalar `ef`. A filling whose boundary lands inside a *degenerate multiplet* — a Hund's-rule atom is exactly this case — where the solve converges and returns an answer 100 % away from a finite difference, because which member falls below the cut is arbitrary; that is the physics, not a missing term. Born effective charges, the dynamical matrix, the strain response, the elastic constants, the Raman tensor and the electrostriction tensors for `nspin = 2`: the solve is spin-polarized, their second-derivative assembly is not. The *screened* response of a magnetic system with vacuum: for `nspin = 2` the kernel `dv_of_drho` is the second derivative of the LSDA energy in the two channel densities and diverges wherever a channel density reaches zero — measured on triplet O<sub>2</sub>, 1504 of 91125 grid points and exactly 1504 NaN. And any potential-only meta-GGA (tb09, bj06), which has no total energy to differentiate — refused by name now rather than surfacing as `v_of_rho` asking for `tau`.

## 8 Optical spectra and excitons: TDDFT

The response above is a *static* one, solved orbital by orbital. An absorption spectrum needs the frequency axis and needs the susceptibility as a **matrix** in reciprocal lattice vectors, so this is the one place where a sum over states earns its keep. `run_absorption` builds

$$\chi_{\mathbf{G}\mathbf{G}'}^0(\omega) = \frac{1}{\Omega} \sum_{\mathbf{k}} \sum_{ij} \frac{(f_i - f_j) \rho_{\mathbf{G}}^{ij} \rho_{\mathbf{G}'}^{ij*}}{\varepsilon_i - \varepsilon_j + \omega + i\eta}, \quad \rho_{\mathbf{G}}^{ij} = \langle u_i | e^{-i\mathbf{G}\cdot\mathbf{r}} | u_j \rangle$$

and solves the Dyson equation with an exchange-correlation kernel:

$$\epsilon^{-1} = 1 + X(1 - X - \tilde{f}_{\text{xc}}X)^{-1}, \quad X = v^{1/2}\chi^0v^{1/2}, \quad \tilde{f}_{\text{xc}} = v^{-1/2}f_{\text{xc}}v^{-1/2}$$



The **bootstrap** kernel of Sharma, Dewhurst, Sanna and Gross (*PRL* **107**, 186401 (2011); Elk's `fxctype = 210`) is that equation's other half, so the two are solved together to self-consistency:

$$f_{xc}^{BS} = -\frac{\epsilon^{-1}(\mathbf{q}, 0) v(\mathbf{q})}{\epsilon_0(\mathbf{q}, 0) - 1}, \quad \epsilon_0 = 1 - v\chi^0$$

It is parameter-free and diverges as  $1/q^2$ , which is what binds an electron-hole pair; ALDA's kernel is finite at  $\mathbf{q} = 0$  where  $v$  diverges, so its head and wings vanish *identically* and no amount of it can bind one.

**Quantum ESPRESSO has no counterpart.** TDDFPT/ is a Liouville-Lanczos solver with RPA and ALDA; it has no bootstrap kernel and no Dyson equation in  $\mathbf{G}$  space. The reference is Elk.

```
from defumat.workflows import run_absorption
import numpy as np

omega = np.arange(0.0, 0.6, 0.005)          # Ry
spectrum = run_absorption(system, pseudos, scf.density, omega,
                          kernel="bootstrap", # or rpa / alda / lrc / bootstrap-1
                          nbnd=60, ecut_response=8.0,
                          broadening=0.012, scissor=0.05)

print(spectrum.absorption)                  # Im eps_M, isotropic average
print(spectrum.frequencies_ev)              # the same grid in eV
print(spectrum.alpha, spectrum.iterations)  # the LRC-equivalent head, and the
                                           # bootstrap's fixed-point count
print(spectrum.static_residual)             # how far the band truncation reaches
```

`chi_0` is the whole cost and the kernel is nearly free, so comparing kernels means building it once: `independent_response` returns it and `solve_dyson` screens it, which is what `run_absorption` does internally and what the notebook shows.

Three knobs here are this code's own. `ecut_response` (Elk's `gmaxrf`) selects the  $\mathbf{G}$ -set the matrix is built over and has its own convergence: too small and the local-field effect is missing, and the cost is quadratic in it. `scissor` is *part of the method* rather than a convenience — the paper's Eq. (3) replaces  $\chi^0$  by a gap-corrected model response, and the velocity matrix elements are renormalised by  $e_{ij}/(e_{ij} \mp \Delta)$  along with it. And `static_residual` is the diagnostic for the one error this phase cannot refuse: how many empty states the sum ran over. It is  $\epsilon_M(0)$  from this run, in RPA, minus the same quantity from the Sternheimer solve, which is band-complete; it is kernel-matched on purpose, since differencing a bootstrap number against a Sternheimer one that contains  $f_{xc}$  would measure the kernel instead.

**A trap that is invisible in the answer:** the macroscopic dielectric function is the inverse of the  $3 \times 3$  *head* of  $\epsilon^{-1}$ , not the head of the inverse of the whole matrix. The second is the microscopic  $\epsilon$ , whose head in RPA is exactly the no-local-field result; it is smooth, positive, has the right peaks and is 9% too large on silicon. Both are returned (`epsilon` and `epsilon_no_local_fields`), because the gap between them is the local-field effect.

Validated by an identity rather than against another code's spectrum: the same  $\epsilon_M(0)$  is reached by this sum over states plus a Dyson inversion and by the projected conjugate-gradient solve of `dielectric_tensor`, which shares no machinery with it and never sees an empty state. On silicon the two agree to **1.3e-2** on a constant of 22 — and that residue is the band truncation, which is why it is reported rather than tuned away. The screening = "hartree" switch on `dielectric_tensor` exists for this comparison: the identity holds only

when both sides use the same kernel.

**Refuses.** Finite  $\mathbf{q}$  (the perturbed states would live at  $\mathbf{k} + \mathbf{q}$ ); ultrasoft and PAW (every matrix element gains the augmentation charge  $Q_{ij}(\mathbf{G})$ ); metals (the  $f_i - f_j$  weight kills the  $i = j$  term, so the intraband Drude response is absent entirely — Elk’s `tdfft1r` does not add it either); `nspin`  $\neq 1$ , noncollinear magnetism, spin-orbit coupling, spin spirals and a Hubbard  $U$ ; a symmetry-reduced  $k$ -set, because symmetrising  $\chi_{\mathbf{G}\mathbf{G}'}^0$  rotates *two*  $\mathbf{G}$  indices at once and the rank- $n$  symmetriser is Cartesian; the `alda` kernel with a gradient-corrected functional, since ALDA’s kernel is pointwise in the density and a GGA’s is not; and **non-convergence of the bootstrap fixed point is an error**, as it is in `tdfft1r.f90`, because a spectrum from a kernel still in motion is the spectrum of nothing.

## 8.1 The conductivity tensor, the Kerr angle and the anomalous Hall effect

The whole complex tensor, interband plus a Drude term:

$$\sigma_{ij}(\omega) = \frac{i}{\Omega} \sum_{\mathbf{k}} \sum_{n,m} \frac{W_n (1 - f_m)}{\epsilon_{mn}} \left[ \frac{v_{nm}^i v_{mn}^j}{\omega - \epsilon_{mn} + i\eta} + \frac{\overline{v_{nm}^i v_{mn}^j}}{\omega + \epsilon_{mn} + i\eta} \right] + \frac{\omega_{p,ij}^2}{4\pi(\gamma - i\omega)}$$

with  $v = \partial H / \partial \mathbf{k}$  from one `jvp` of  $H(\mathbf{k})$  at a frozen sphere. *That* is the one thing here not transcribed from `dielectric.f90`, and it is not a refinement: `epsilon.x` accumulates  $\langle \psi_1 | \mathbf{G} | \psi_2 \rangle$ , and  $[H, \mathbf{r}] \neq \mathbf{p}$  when the pseudopotential is nonlocal.

What is new beside `run_absorption` is the **antisymmetric** part,  $\sigma_{xy}$  — the response that rotates light reflected off a magnet, and whose static limit is the intrinsic anomalous Hall conductivity. Time reversal forces it to zero, so a nonmagnetic crystal has none however strong its spin-orbit coupling; a *collinear or coplanar* magnet without spin-orbit coupling has none either, because a spin rotation composed with conjugation survives as an antiunitary symmetry. (A *noncoplanar* magnet does have one with no spin-orbit coupling at all — the topological Hall effect — so the usual pairing is a route and not a theorem.)

```
calc = Calculator.from_file("ni-soc.in", pseudo_dir="pseudo")

sigma = calc.get_optical_conductivity(nbnd=36, window=0.6, nw=200)
print(sigma.sigma_s_per_cm.shape)    # (nw, 3, 3) in S/cm
print(sigma.kerr)                    # complex Kerr angle, degrees, per frequency
print(sigma.hall_conductivity)       # the static antisymmetric part, S/cm
print(sigma.plasma_ev)               # hbar omega_p, eV, a (3, 3) tensor
print(sigma.dielectric)              # 1 + 4 pi i sigma / omega
print(sigma.fsum, sigma.band_cut_gap) # the two diagnostics to read
```

There is no reference output for any of it, so the validation is internal and analytic. The *symmetric* part reproduces the independent-particle dielectric function of the TDDFT stack to **3.6e-14**, which is a different assembly (a Dyson solve over a response sphere) reaching `ph.x` through `dielectric_tensor`. The **f-sum rule** is the absolute anchor:  $\int_0^\infty \text{Re } \sigma_{aa} d\omega$  must equal  $(\pi/2\Omega) \sum W_n \langle n | \partial^2 H / \partial k_a^2 | n \rangle$ , and the familiar  $\pi n_e / 2$  is only the *local* Hamiltonian’s special case — silicon’s diamagnetic weight is measurably **0.943** and not one. Measured against a central difference of the velocity operator, the ratio is 1.258 on a  $4^3$  grid, 1.045 on  $6^3$  and **0.990** on  $8^3$ : it converges in the *k-grid*, because what separates the two is  $\sum_k w_k \partial^2 \epsilon_n / \partial k^2$ , which vanishes over the zone by periodicity and not on

a coarse mesh. Aluminium's plasma frequency comes out at **12.98 eV** against a free-electron 16.27, which is the zone-boundary gaps removing Fermi surface, and its tensor is isotropic to **1.7e-4 eV** with nothing imposing that.

**Refuses. Ultrasoft and PAW**, for the Kubo curvature's reason: with a moving  $\hat{S}$  the current operator carries  $\epsilon_n \partial S / \partial k$  off the diagonal, and that term is identically zero for a norm-conserving dataset, so nothing validated here can see whether its convention is right. A **spin spiral**, whose two spinor components live on different spheres. A **symmetry-reduced k-set**, because the antisymmetric part is an axial vector — the escape is the whole unshifted grid, which is closed under the point group. And **nspin = 2**, whose two channels are two band structures whose conductivities add; a spinor run is the regime a magneto-optical spectrum wants anyway, since  $\sigma_{xy}$  needs spin-orbit coupling and collinear spin has none.

**Two things are reported rather than refused, and both must be read.** `band_cut_gap` is the gap between the last band summed and the first dropped: stopping *inside* a degenerate multiplet keeps some members and drops others, and silicon's antisymmetric  $\sigma$  — zero by time reversal — comes out at 4.0e-13 when the cut falls at a real gap and at 1.0e-5 when it does not. And the **anomalous Hall conductivity of a metal is mesh-hungry**: the static Berry-curvature integral is concentrated on near-degeneracies at the Fermi surface, so it is quoted with its  $k$ -convergence and not on its own.

## 8.2 The shift current: the bulk photovoltaic effect

Illuminate a crystal with no inversion centre and it carries a direct current, with no junction and no built-in field. The intrinsic part of that is the *shift current*: the photoexcited electron is born displaced, and what the current counts is the real-space shift between the valence and conduction Wannier centres. It is a second-order optical response,  $J^a = 2\sigma^{abc}(0; \omega, -\omega)E_b(\omega)E_c(-\omega)$ , with

$$\sigma^{abc}(\omega) = -\frac{i\pi e^3}{4\hbar^2} \int [d\mathbf{k}] \sum_{n,m} f_{nm} \left( I_{mn}^{abc} + I_{mn}^{acb} \right) [\delta(\omega_{mn} - \omega) + \delta(\omega_{nm} - \omega)], \quad I_{mn}^{abc} = r_{mn}^b r_{nm}^{c;a}.$$

The whole difficulty is  $r_{nm}^{c;a}$ , the *generalised derivative*  $\partial_a r_{nm}^c - i(A_{nn}^a - A_{mm}^a)r_{nm}^c$  — neither piece of which is separately computable, since the diagonal Berry connection is gauge dependent and is not a function of  $H$  at one  $\mathbf{k}$ . The escape is the Aversa-Sipe sum rule, which writes the covariant object entirely in gauge-free quantities living at a single  $\mathbf{k}$ :

$$r_{nm}^{c;a} = \frac{i}{\omega_{nm}} \left[ \frac{v_{nm}^c \Delta_{nm}^a + v_{nm}^a \Delta_{nm}^c}{\omega_{nm}} - w_{nm}^{ca} + \sum_{p \neq n,m} \left( \frac{v_{np}^c v_{pm}^a}{\omega_{pm}} - \frac{v_{np}^a v_{pm}^c}{\omega_{np}} \right) \right],$$

with  $v_{nm}^a = \langle n | \partial_a H | m \rangle$  and  $w_{nm}^{ab} = \langle n | \partial_a \partial_b H | m \rangle$ . Both are derivatives of code that already exists:  $v$  is one jvp of  $H(\mathbf{k})$  at a frozen sphere and  $w$  is that jvp differentiated by a second one. Nothing is transcribed. The interband dipole  $r_{nm}^a = -i v_{nm}^a / \omega_{nm}$  is *exact* for an eigenstate of the full  $H(\mathbf{k})$ , which is why a plane-wave code needs no counterpart to Wannier90's position-operator matrices.

```
calc = Calculator.from_file("alas-raman.in", pseudo_dir="pseudo")

sc = calc.get_shift_current(nbnd=14, window=0.9, nw=180)
print(sc.sigma.shape)           # (nw, 3, 3, 3) in A/V^2, real
```

```

print(sc.component(0, 1, 2))    # sigma^xyz(omega), -43m's only free component
print(sc.frequencies_ev)      # the photon energy axis in eV
print(sc.truncation)           # the band-truncation diagnostic: read it
print(sc.band_cut_gap)         # where the band set was cut, in Ry

```

There is no reference binary — `pw.x` computes no photocurrent of any kind, and Elk's `nonlinopt.f90` is second-harmonic generation, a different response — so the validation is internal and it is three statements. The **sum rule against a parallel-transport finite difference** of the dipole, which shares nothing with it: no second derivative, no intermediate sum, only the phases of neighbouring **k**-points fixed to one another. That comparison is decisive in exactly one form, because the sum rule is an *identity over a complete basis*: held on a frozen sphere of 158 plane waves, AlAs agrees to **1.8e-4** at 158 bands against 4.3e-2 at 120 and 6.0e-2 at 20. A residue that falls off a cliff at completeness is a correct assembly with a truncated sum; one that plateaus is a bug. Then the **crystal class**: AlAs comes out exactly  $\bar{4}3m$  on a `nosym` grid, only the components with all three labels different surviving, while centrosymmetric silicon gives zero from the same machinery. And the **absorption edge**:  $\sigma^{xyz}$  is 6.9e-25 A/V<sup>2</sup> at 2 eV against a peak of 1.1e-4 near 4.4 eV, which is the one check no index bookkeeping can pass by accident — no absorption, no photocurrent.

**All three are blind to an overall constant**, and this document says so rather than leaving it to be discovered: a  $\sigma$  wrong by a factor of two is still exactly  $\bar{4}3m$ , still zero on silicon and still zero below the gap. Two further checks close that gap. The **spin sum** — the same cell run `nspin = 1` and as a two-component spinor gives the same tensor to **4.6e-9**, peak ratio 1.000000. And the **overall scale against the published literature**, where the *ordering* is what makes the comparison sharp rather than merely consistent. AlAs's first peak converges to **35.0**  $\mu\text{A/V}^2$  at 4.17 eV (33.9 at `ecutwfc = 16`, then 33.7, 34.8 and 35.0 at 22 with 14, 22 and 30 bands), where calculations across the fourteen III-V and II-VI zincblende semiconductors span 14  $\mu\text{A/V}^2$  (CdSe, smallest) to 83 (AlSb, largest) and find the aluminium compounds the strongest of the family. A factor of two either way breaks that ordering rather than moving the number: 17 would put AlAs at the bottom of the family, below the cadmium chalcogenides, and 70 would put it beside AlSb, the heaviest and narrowest-gap member. The number to compare is the peak below about 6 eV, which is the range published spectra cover. That sweep also showed `truncation` to be *predictive* rather than merely indicative: it falls 3.5% to 1.0% over the band sweep while the value moves 3.9%.

**The convention is declared, because the literature has two.** What is implemented is  $j^a = 2\sigma^{abc}(0; \omega, -\omega) \text{Re}[\mathcal{E}_b(\omega)\mathcal{E}_c(-\omega)]$ , the definition of the paper the formulae above come from. Some authors write  $J = \sigma\mathcal{E}\mathcal{E}$  and report numbers twice as large for the same physics, so a comparison against a paper that does not state its normalisation is worth nothing to better than a factor of two.

**Refuses. Ultrasoft and PAW**, for the Kubo curvature's reason one order further out: with a moving  $\hat{S}$  the dipole carries  $\partial S/\partial k$  and its derivative carries  $\partial^2 S/\partial k_a \partial k_b$ , both identically zero for a norm-conserving dataset, so nothing validated here can see whether their convention is right. A **spin spiral**. A **symmetry-reduced k-set**, because  $\sigma^{abc}$  is a polar rank-3 tensor and a wedge sum is not the cell's until it is averaged over the point group; the escape is the whole unshifted grid. A **metal**, whose partially filled bands contribute pairs with vanishing  $\omega_{nm}$  where every denominator above diverges — the Fermi-surface terms that replace them are a different assembly. **DFT+U**, whose Hubbard term is built from  $\phi_U(\mathbf{k} + \mathbf{G})$  and so carries a velocity and a

second velocity of its own. And `nspin = 2`. A *spinor* run is supported and tested. **The band count is reported rather than refused, and it is the largest approximation here.** The intermediate sum over  $p$  converges only as it approaches completeness, so a production band count carries a few per cent that no symmetry check sees — truncation is the in-run estimate and is taken over the top *quarter* of the bands, because dropping one band reports 5.3e-5 where the error is 4 per cent. The route that would remove it is written down in `NONLINEAR.md`: the two sums are a Sternheimer solve rather than a spectral sum, and what stops it being taken is that the operator to invert is indefinite. `band_cut_gap` carries the same warning it does for the conductivity.

### 8.3 Second-harmonic generation

Shine light of frequency  $\omega$  on a crystal with no inversion centre and some of it comes back out at  $2\omega$ . The tensor that says how much is the second-order optical susceptibility  $P^a(2\omega) = \chi^{abc}(-2\omega; \omega, \omega) E_b(\omega) E_c(\omega)$ , a polar rank-3 tensor, symmetric in its last two labels and identically zero in every centrosymmetric crystal. It is a sum of three physically distinct pieces — the pure interband term  $\chi_{II}$ , the intraband modulation of it  $\eta_{II}$ , and the modulation of the transition energy  $\frac{i}{2\omega}\sigma_{II}$  — each of which is a contraction of the interband dipoles  $r_{nm}^a = -i v_{nm}^a / \omega_{nm}$  over two resonant denominators, one at  $\omega$  and one at  $2\omega$ :

$$\chi^{abc}(-2\omega; \omega, \omega) = \frac{1}{\Omega} \sum_{\mathbf{k}, m, n} \left[ \frac{c_{mn}^{(1)}}{\omega_{mn} - \omega + i\eta} + \frac{c_{mn}^{(2)}}{\omega_{mn} - 2(\omega - i\eta)} \right],$$

where the coefficients  $c_{mn}^{(1,2)}$  carry a further sum over an intermediate state and are built once, before the frequency axis is touched. This is the same sum over states the shift current uses, contracted differently; it needs no second derivative of  $H$  at all, because that triple sum is the sum-rule expansion of the generalised derivative written out rather than collapsed.

```
calc = Calculator.from_file("alas-shg.in", pseudo_dir="pseudo")

shg = calc.get_shg(nbnd=22, window=0.6, nw=240, broadening=0.010)
print(shg.chi.shape)           # (240, 3, 3, 3), complex, in pm/V
print(shg.component(0, 1, 2)[0]) # chi^xyz at the lowest frequency
print(shg.d_coefficient(0, 1, 2)[0]) # d = chi/2, what the literature quotes
print(shg.frequencies_ev[:3])    # the FUNDAMENTAL photon energy, in eV
print(shg.truncation)            # the band-truncation diagnostic
print(shg.chi_ii, shg.eta_ii, shg.sigma_ii) # the three parts, kept apart
```

**There is a real reference here**, which is unusual in this corner of the package: Elk's `nonlinopt.f90` (task 125) computes the same tensor, and `tests/data/elk/` carries its output for the same ALAs crystal, the three parts separately as well as their sum. `pw.x` has none — its `el_opt.f90` is the *static* electro-optic response, and the branch that reaches even that is the one this document's Raman section records as broken in the vendored 7.5.

What the comparison is *not* is digit-for-digit, and saying so is part of quoting it: an all-electron LAPW basis is not a norm-conserving pseudopotential one, and a second-order susceptibility carries two energy denominators, so a gap difference is not a scale factor. On the same crystal and the same  $6^3$  mesh: the **resonance position** agrees to **0.5%** (2.152 eV against 2.163), the **peak height** to **7%** (27.52 a.u. against 25.70), and the **static value** to **11%** (−3.10 a.u. against −3.50). The last is the one that moves with the

basis, and it is shown converged: `ecutwfc` 10, 30 and 45 give  $-2.71$ ,  $-3.10$  and  $-3.13$ , so the residue is the pseudopotential rather than the cutoff.

Four internal checks close what the reference cannot. `43m`: `AlAs` comes out exactly zincblende on a `nosym` grid with nothing imposing it. **Silicon**: an inversion centre forbids every component, and the centrosymmetric control runs through the same machinery. **The two-photon edge**:  $\text{Im } \chi$  turns on at *half* the smallest direct gap, not at the gap, because the second channel's denominator is  $\omega_{mn} - 2\omega$  — the only check here blind to both the scale and the symmetry, and a lost factor of two there moves the absorption edge by an octave. And **the spinor**, which is the scale's own anchor: the same cell run `nspin` = 1 and as a two-component spinor must give the same tensor, which is what catches a factor of two in the spin sum.

**The convention is declared, because the literature has two.** What is computed is  $\chi^{(2)}$ , defined by  $P^a = \chi^{abc} E_b E_c$ ; nonlinear-optics papers more often quote  $d^{abc} = \chi^{abc}/2$ , so a number compared against a published  $d$  without halving it is wrong by exactly two and nothing about its shape, symmetry or zeros would say so. `d_coefficient` returns the halved form.

**The scissors shift is the one knob worth explaining.** An LDA or GGA gap error does not scale a second-order susceptibility, it moves its resonances, and `scissor` applies the usual rigid shift to the empty states. The dipoles are built from the *unshifted* gaps — a scissors correction is a statement about energies, not about wavefunctions — which is what Elk reaches by scaling its momentum matrix elements instead. Validated against Elk's own scissored run: at  $\Delta = 0.05$  Ha the  $2\omega$  peak moves **0.0502** Ry against a half-scissor of 0.0500, and its height falls to **0.60** of the unscissored value against Elk's 0.58.

**Refuses.** The same five the shift current refuses, since this is the same velocity matrix elements contracted a different way and it inherits every reason: **ultrasoft and PAW** (a moving  $\hat{S}$  puts  $\partial S/\partial k$  in the dipole), a **spin spiral**, a **symmetry-reduced  $k$ -set** (a polar rank-3 tensor is not the cell's until averaged over the point group; the escape is the whole unshifted grid), a **metal, DFT+ $U$** , and `nspin` = 2. A *spinor* run is supported and tested.

**This does not lift the  $\chi^{(2)}$  refusal in the Raman section**, and the distinction matters. That one is about the  $2n + 1$  route through the Sternheimer solver, where a field enters only through a source term and the  $\langle u_i | r_k | u_j \rangle$  piece has nothing to build it from; a sum over states never needed that object. The electro-optic tensor and a truncation-free static  $\chi^{(2)}$  are still refused for the same missing term.

**The band count is reported rather than refused**, exactly as for the shift current and for the same reason: the sum over the intermediate state is an identity only over a complete basis. `truncation` is the in-run estimate, taken over the top quarter of the bands.

## 9 Topological invariants

Everything is built from one primitive — the overlap of occupied manifolds at neighbouring  $k$ -points,

$$M_{mn}^{(\mathbf{b})} = \langle u_{m\mathbf{k}} | \hat{S} | u_{n,\mathbf{k}+\mathbf{b}} \rangle$$

because a determinant of overlaps is blind to the unitary mixing a degenerate eigensolver leaves. The Berry curvature is the phase of a plaquette of them and the Chern number



their lattice sum, which is an **exact integer on any mesh**:

$$F_{12}(\mathbf{k}) = \ln [U_1(\mathbf{k})U_2(\mathbf{k}+\hat{1})U_1(\mathbf{k}+\hat{2})^{-1}U_2(\mathbf{k})^{-1}], \quad C = \frac{1}{2\pi i} \sum_{\mathbf{k}} F_{12}(\mathbf{k})$$

$Z_2$  has two independent routes — the flow of Wannier charge centres, and, when there is an inversion centre, the parity product over the eight TRIM points:

$$(-1)^\nu = \prod_{i=1}^8 \prod_{n \text{ occ}}^{\text{Kramers}} \xi_{2n}(\Gamma_i)$$

```
from defumat.workflows.topology import (run_z2, run_z2_3d,
                                         run_berry_curvature)

z2 = run_z2(system, pseudos, scf.density, axis=2) # Wilson loop by default
print(z2.z2, z2.gap_step) # read gap_step before believing the integer

chern = run_berry_curvature(system, pseudos, scf.density, shape=(12, 12))
print(chern) # an exact integer on any mesh

nu = run_z2_3d(system, pseudos, scf.density) # the four 3D indices
print(nu) # (nu0; nu1 nu2 nu3)
```

**Refuses.** A collinear spin-polarized run (two Hamiltonians, so two independent sets of invariants and no statement of which is meant); a  $Z_2$  without spin-orbit coupling, where every invariant is zero for a reason that has nothing to do with the band structure; the parity route without an inversion centre; and a manifold that is not separated from the band above it (`gap_tol`). A PAW run needs the converged `becsum` passed beside the density — `run_z2(..., becsum = scf.becsum)`, which calculator does for you — because `ddd_paw` is a property of the wavefunctions and cannot be rebuilt from a density; leaving it out is worth 1.03 eV in the eigenvalues, and an invariant computed from those is still an integer.

**Running both  $Z_2$  routes is the check.** Where they disagree the parity one is the answer, because it has no mesh and the Wilson route does. **Two things bite in a plane-wave code and both are silent:** neighbouring  $k$ -points do not share a  $G$ -sphere, so coefficients are aligned by Miller index; and the wrap at the zone edge is a *shift* of that index, without which the Chern number comes out smooth and non-integer.

## 9.1 The curvature as a map: the Kubo route

Two constructions of the same quantity, chosen by name:

$$\Omega_n^{12}(\mathbf{k}) = -2 \operatorname{Im} \sum_{m \neq n} \frac{A_{nm}^1 A_{mn}^2}{(\varepsilon_n - \varepsilon_m)^2}, \quad A_{nm}^a = \langle \psi_n | \partial_{k_a} H - \varepsilon_n \partial_{k_a} S | \psi_m \rangle$$

with `method="kubo"`, against the lattice field strength of Fukui, Hatsugai and Suzuki with `method="fhs"` (the default). The velocity operator is one jvp of  $H(\mathbf{k})$  at a frozen plane-wave sphere — no dense  $H(\mathbf{k})$  is ever formed — and the tangent for a crystal direction  $d$  is the reciprocal lattice vector  $\mathbf{b}_d$ .



```

from defumat.workflows.topology import run_berry_curvature

curvature = run_berry_curvature(
    system, pseudos, density, shape=(4, 4), nocc=4, nbnd=12, method="kubo"
)
curvature.curvature          # (4, 4),  $\Omega(k)$ 
curvature.curvature_by_band  # (4, 4, nocc) -- see the refusal box
curvature.truncation         # what the highest empty band was worth

```

Validated against the FHS flux, which shares no machinery with it: a centred plaquette shrunk around one  $\mathbf{k}$ -point of zincblende AIs converges onto the pointwise Kubo value to  $3.2 \times 10^{-3}$ , and the whole  $24 \times 24$  mesh agrees plaquette by plaquette to  $1.45 \times 10^{-4}$ , improving 65-fold over a sixfold refinement. The velocity matrix elements themselves reproduce a central finite difference of  $H(\mathbf{k})$  to  $1.8 \times 10^{-9}$ , and on centrosymmetric silicon the curvature vanishes *pointwise* to  $3.5 \times 10^{-5}$ .

**Refuses.** `method="kubo"` is for the map, never for the integer: its denominator  $(\varepsilon_n - \varepsilon_m)^{-2}$  is singular at a degeneracy and its Brillouin-zone sum is an ordinary Riemann sum, which converges to an integer without ever being one — `method="fhs"` is the default and is exact on any mesh. **Ultrasoft and PAW are refused by name:** the  $\varepsilon_n \partial_k S$  term is identically zero for a norm-conserving dataset, so no validation here can see whether its off-diagonal convention is right, and  $q_{ij}(\mathbf{b})$  is the second missing piece; `method="fhs"` carries both correctly on all three dataset kinds. The sum over empty states is truncated at `nbnd` and the truncation is *reported* rather than hidden — read `BerryCurvature.truncation`, or `truncation_abs` wherever symmetry forces the curvature to vanish, since the ratio is then noise over noise. `curvature_by_band` is gauge invariant only for a non-degenerate band; inside a degenerate multiplet only the sum over the members is defined, and `curvature` (the manifold total) always is.

## 9.2 The Berry-phase polarization

The polarization of a crystal is not the dipole of its unit cell — that integral depends on where the cell is cut. What is a property of the crystal is a *phase*: the Berry phase the occupied manifold accumulates as  $\mathbf{k}$  crosses the zone once along a reciprocal lattice vector  $\mathbf{b}_i$  (King-Smith and Vanderbilt, PRB **47**, 1651 (1993)),

$$\Omega \mathbf{P} = \sum_i \frac{\varphi_i}{2\pi} \mathbf{a}_i, \quad \varphi_i = \underbrace{\text{Im} \ln \prod_j \det M^{(\mathbf{b}_i/N)}(\mathbf{k}_j)}_{\text{electronic}} + \underbrace{2\pi \sum_a Z_a^v s_{ai}}_{\text{ionic}}$$

with  $M$  the same overlap the invariants above are built from and  $s_{ai}$  the atom's crystal coordinate. It is the *determinant* of the overlap that is kept here where a Wilson loop keeps the eigenvalues, which is the whole difference between the two calculations.

**The value is defined modulo a quantum** and the result object carries it. A single number is not a physical statement: what is meaningful is a *difference* between two geometries taken on the same branch, which is what a Born effective charge and a piezoelectric constant are made of.

```

from defumat.workflows.polarization import run_polarization

p = run_polarization(system, pseudos, scf.density,

```

```

        gdir=2,          # 0-based; pw.x's gdir = 3
        nppstr=6,       # as pw.x counts it
        transverse=(2, 2))
print(p.total_phase, p.quantum)  # -0.2412383 (mod 1)
print(p.polarization_si)        # -0.34909 C/m^2, mod 1.44708
print(p.strings.phases)         # one phase per string, four here

```

`Calculator.get_polarization()` reads `lberry`, `gdir` and `nppstr` straight off a `pw.x` input, so a polarization run transfers unchanged. `nppstr` counts the repeated endpoint the way `kp_strings` does; the mesh built here carries `nppstr - 1` distinct points and closes the string with the reciprocal lattice vector, which is the same links and one diagonalisation fewer per string.

**Refuses. A metal** — a Berry phase is a property of a gapped manifold, and a smeared occupation does not say which bands the string carries. `nspin = 2`: two channels are two independent string sets and one phase each, which is a layout this does not have rather than a missing term. And a **spin spiral**, whose two spinor components live on spheres centred at  $\mathbf{k} \pm \mathbf{q}/2$ , so a string overlap is not a single gather. A *spinor* run (`noncolin`) is supported and has its own `pw.x` reference; its quantum is half the scalar one, because a spinor band holds one electron. Ultrasoft and PAW are *not* refused, and both are validated: zincblende SiC matches `pw.x` string by string with an electronic phase of  $-0.009696$  against  $-0.00970$ . It had to be SiC rather than silicon, because a centrosymmetric crystal's electronic phase is zero whatever  $q_{ij}(\mathbf{b})$  does and so agrees with zero however wrong the augmentation is.

**Validated three ways, two of which need no other code.** Against `pw.x`'s own `lberry` run on AlAs it agrees *string by string* and on every digit printed (electronic phase 0.00876, total  $-0.24124$ ). The Zak phase of the SSH model is pinned to 0 or  $\pi$  by symmetry and comes out exactly that on a mesh of *any* size ( $10^{-10}$ ), and silicon is pinned by inversion to 0 or half a quantum, landing on the half ( $10^{-8}$ ). Silicon is also the second `pw.x` reference, and the one that exercises the quantum of 2: both its species carry an even valence, where AlAs's arsenic does not. The strongest is  $dP/du$  against the Born charges, which come from a Sternheimer solve and share no machinery with a string of overlaps: AlAs on an  $8 \times 8 \times 8$  ground state gives  $Z^*(\text{Al}) = +2.148$  against `ph.x`'s  $+2.142$ , charge-neutral to  $5 \times 10^{-5}$ .

### 9.3 The magnetoelectric tensor

How much electric polarization a magnetic field induces,

$$\alpha_{ij} = \frac{\partial P_i}{\partial B_j},$$

computed Elk's way (task 390): a full ground state at  $B \mp \delta/2$  for each Cartesian direction, the Berry-phase polarization of each, and a central difference. The cheap route — one derivative of the magnetization along an electric field's response — needs a noncollinear Sternheimer solve, and an electric field as a ground-state perturbation, which this package does not have.

**Two symmetries must both be broken or the answer is zero for a reason that has nothing to do with the calculation:** time reversal and inversion each map  $\alpha$  to  $-\alpha$ . Elk's own example is built around the first, and so is the case here — a non-magnetic semiconductor has no linear tensor at all, so a *large* external field is applied to break time reversal by hand and the derivative taken against a small further change in it. The

magnetism is induced rather than spontaneous, which is what makes a plain gapped semiconductor usable.

```
from defumat import Calculator

alas = Calculator.from_file("alas-magnetoelectric.in", pseudo_dir)
me = alas.get_magnetoelectric_tensor(delta=0.02, directions=(2,))
print(me.magnitude)          # 4.4376e-06
print(me.largest_step)       # how close a field step came to a branch
```

**Refuses.** An unconverged ground state, which is what makes the next limitation visible rather than silent. **Only the column parallel to the seeded magnetization is reachable** on the committed cell: a field along the direction `starting_magnetization` points converges in 12 iterations, and one with a transverse component does not converge in 150, because rotating the moment is the slow mode. It is not the tilt — a pure transverse field fails on its own. Seed the magnetization along the column you want. A run carrying no magnetization vector ( $n_{\text{spin,mag}} \neq 4$ ): the Zeeman term then multiplies nothing and every polarization comes out identical, which looks like a symmetry result. And a field step that moves the Berry phase by more than a quarter of its quantum, since a difference taken across a branch is the quantum divided by  $\delta$  — large, smooth and meaningless. Everything the polarization refuses is refused here too.

**The validation is the spin-orbit null.** Without the coupling a global spin rotation maps  $B$  to  $-B$  and leaves every charge observable fixed, so  $\alpha$  vanishes identically. On the same cell, same field, same ultrasoft family at the same valence, with only `lspinorb` and the scalar/fully relativistic datasets changed:  $3.43 \times 10^{-9}$  against the coupled  $4.44 \times 10^{-6}$ , a ratio of **1293**. Two things had to be right together to get there, and neither helps alone: `chain` (seeding one field run from the other, which breaks the very  $B \rightarrow -B$  symmetry the null rests on — off by default here, unlike Elk) and the *polarization's* threshold rather than the self-consistent field's. **The absolute value is not calibrated against another code:** Elk's `bfieldc` and this package's `B_field` differ in normalisation by about two orders of magnitude, and Elk's own tensor on this cell is not converged.

## 10 Effective masses, site angular momenta, nesting and structure factors

Quantities Elk computes and `pw.x` does not. The first two cost one NSCF at a single  $k$ -point; the nesting function costs one NSCF on a dense grid, and nothing after it; the structure factors cost one transform of a density that already exists.

### 10.1 The effective mass tensor

$$\left(\frac{1}{m^*}\right)_{ab} = \frac{1}{2} \frac{\partial^2 \varepsilon_n(\mathbf{k})}{\partial k_a \partial k_b},$$

in units of  $1/m_e$ . The factor of a half is Rydberg atomic units and not a normalisation:  $\hbar^2/2m_e$  is exactly 1 Ry bohr<sup>2</sup>, so a free electron has  $\varepsilon = |\mathbf{k}|^2$  and the tensor is the identity.

**The first derivative is analytic and only the second is a difference.**  $\partial \varepsilon_n / \partial k$  is the generalised Hellmann–Feynman expression built from one jvp of  $H(\mathbf{k})$ ; its second deriva-

tive is not reachable the same way, because an individual band's  $|\partial_k \psi_n\rangle$  is not what the Sternheimer projector produces (it removes the whole occupied manifold, and the band whose mass is wanted is usually empty). So the construction is one central difference of an *exact* first derivative — six stencil points against Elk's twenty-seven, and no polynomial fit. Elk's route, differencing eigenvalues, is kept beside it as `method="eigenvalue"` because it shares nothing with the velocity operator.

```
from defumat import Calculator

calc = Calculator.from_file("si.scf.in")
mass = calc.get_effective_mass((0.0, 0.0, 0.0), nbnd=10) # crystal coordinates

mass.inverse_mass      # (nbnd, 3, 3) in 1/m_e, cartesian
mass.principal(band=7) # the three principal masses and their axes
mass.density_of_states_mass(band=7) # (m_x m_y m_z)^(1/3)
mass.truncation         # the O(h^2) the Richardson step removed
mass.multiplets         # per multiplet, including the degenerate ones
```

On two-atom silicon at  $\Gamma$  against the vendored **Elk** binary (all-electron LAPW, PBE,  $a = 10.26$  bohr) with a PAW dataset here: the non-degenerate  $\Gamma_{1v}$  agrees to **0.02%** (0.8600902 against 0.8603044, i.e.  $m^* = 1.16267 m_e$  against 1.16238) and  $\Gamma_{2'c}$  to 0.36% ( $m^* = 0.170439 m_e$ ). The two routes here agree with each other to  $1.2 \times 10^{-5}$ , and the tensors come out isotropic to  $2.9 \times 10^{-8}$  with nothing imposing cubic symmetry.

**Refuses.** A band inside a **degenerate multiplet** has no tensor of its own — the eigensolver's arbitrary rotation within the manifold rotates it — so those bands come back as nan and what is reported instead is the multiplet's *summed* inverse mass, which is the trace and is invariant. Silicon's threefold  $\Gamma_{25'}$  is the case, and Elk's own output is the evidence: it prints  $-5.8938, -3.7690, -3.7690$  for three states whose energies agree to  $5 \times 10^{-8}$  Ha. A **spin spiral** is refused, since a k-stencil moves both spheres. `delta` is the stencil's outermost displacement in both routes; the  $O(h^2)$  truncation is removed by a Richardson step and then *reported* in `truncation` rather than tuned away.

## 10.2 $\langle L \rangle$ , $\langle S \rangle$ and $\langle J \rangle$ on each atom

Where the orbital moment of a spin-orbit magnet actually sits. From one site density matrix per atom and shell, built out of the same projection a projected density of states uses,

$$\rho_{(ms),(m's')}^a = \sum_{n\mathbf{k}} w_{n\mathbf{k}} c_{ms,n\mathbf{k}}^a \overline{c_{m's',n\mathbf{k}}^a}, \quad c = \langle \phi | S | \psi \rangle,$$

and then  $\langle L_i \rangle = \sum_s \text{Tr}_m(L_i \rho_{ss})$  and  $\langle S_i \rangle = \frac{1}{2} \sum_m \text{Tr}_s(\sigma_i \rho_{mm})$ , with  $J = L + S$ . The  $L$  matrices are the complex-basis ones conjugated into the *real* spherical harmonics the code works in, using the same `rot_ylm` unitary the spin-orbit projectors are built from.

```
from defumat import Calculator

calc = Calculator.from_file("ni.scf.in") # noncolin, lspinorb, nosym
momenta = calc.get_angular_momenta()

print(momenta.table()) # what Elk writes to LSJ.OUT
momenta.atoms[0].l     # <L> on atom 1, cartesian, units of hbar
```

```
momenta.atoms[0].j          # <J> = <L> + <S>
momenta.total_orbital       # summed over the cell
```

**Validated by identities, not by another code’s number** — Elk integrates over a muffin tin of a stated radius and this projects onto an orbital set, so the two differ by definition the way Löwdin and muffin-tin charges do.  $\langle L \rangle$  is *quenched to zero* without spin-orbit coupling, measured at  $1.7 \times 10^{-16}$  on silicon; with the coupling on, fully-relativistic nickel gives  $\langle L_z \rangle = 0.0364767 \hbar$  against a measured orbital moment of about  $0.05 \mu_B$ , with  $|L|/|S| = 0.11665$  against an experimental  $m_L/m_S \approx 0.1$  and  $\langle L \rangle \parallel \langle S \rangle$  (Hund’s third rule). Driving the moment along  $z$ ,  $x$  and  $y$  gives  $|\langle L \rangle| = 0.0364767$  in all three, with a spread of  $7.3\text{e-}11$  over the three cubic axes, with nothing imposing that a magnitude is a scalar.

**Refuses.** A **symmetry-reduced k-set**:  $\langle L \rangle$  and  $\langle S \rangle$  are vectors, and summing a wedge sums a wedge — the axial-vector group average with  $\det(R)$  and the time-reversal sign is not written here. Run the whole grid with `nosym = .true.`, unshifted so that it is closed under the point group. A **fully-relativistic ultrasoft or PAW** dataset, because the spinor overlap’s off-diagonal spin blocks are `qq_so` and the projection here applies the scalar  $S$  to each component; a fully-relativistic *norm-conserving* dataset has  $S = 1$  and is exact. And a **spin spiral**, whose two components do not share a set of atomic orbitals. `pw.x`’s `lorbm` computes the *cell*’s orbital magnetization and has no per-atom counterpart to check this against.

### 10.3 The Fermi-surface nesting function

How much of the Fermi surface maps onto itself when translated by  $\mathbf{q}$ :

$$N(\mathbf{q}) = \frac{1}{N_k} \sum_{\mathbf{k}} g(\mathbf{k}) g(\mathbf{k} + \mathbf{q}), \quad g(\mathbf{k}) = g_s \sum_{s,n} \delta(\varepsilon_{sn\mathbf{k}} - E_F).$$

$g(\mathbf{k})$  is the density of states *at one k-point* — a picture of the Fermi surface on the grid — so  $N(\mathbf{q})$  is the  $\omega \rightarrow 0$  imaginary part of the Lindhard function stripped of its matrix elements: the *geometric* half of an instability. Where it is large, a perturbation of wavevector  $\mathbf{q}$  connects many occupied states to many empty ones at no energy cost, which is what softens a phonon, opens a charge-density-wave gap, or picks a spin spiral’s pitch.

**It is a convolution.** Elk writes an  $O(N_q N_k)$  double loop with  $\mathbf{k} + \mathbf{q}$  folded back onto the grid; that fold makes the sum a cyclic cross-correlation, so  $N = \text{ifftn}(|\text{fftn}(g)|^2)/N_k$  gives every  $\mathbf{q}$  in one transform. Measured on a  $12^3$  aluminium grid that is **0.001 s** here against Elk’s 0.37 s. A symmetry-reduced wedge is *unfolded* onto the complete grid rather than refused, since  $\varepsilon_n(R\mathbf{k}) = \varepsilon_n(\mathbf{k})$ : 72 diagonalisations instead of 1728.

```
from defumat import Calculator

calc = Calculator.from_file("al.scf.in")
nest = calc.get_nesting(grid=(12, 12, 12)) # the grid is the convergence knob

nest.nesting      # (nq,) N(q), in (states/Ry/cell)^2
nest.qpoints      # (nq, 3) crystal coordinates, on the unshifted grid
nest.ratio        # N(q) / D(E_F)^2, dimensionless and 1 on average
nest.peak()       # the largest N(q) away from q = 0, and where it is
nest.along(2)     # (q_3, N) along one crystal axis, for a line plot
nest.as_grid()    # (n1, n2, n3), for a slice
nest.fermi_dos    # D(E_F) in states/Ry/cell
```

```
nest.elk_units      # the same array in NEST3D.OUT's normalisation
```

The functional entry point is `defumat.workflows.run_nesting`, which takes the converged density directly.

**Two identities come with the normalisation and both are worth reading.** The mean of  $N$  over the whole  $\mathbf{q}$ -grid is exactly  $D(E_F)^2$  (`sum_rule` reports the residue; it is  $3 \times 10^{-16}$  relative). And  $N(0) \geq N(\mathbf{q})$  for every  $\mathbf{q}$  by Cauchy-Schwarz, so **the nesting peak is always a peak away from the origin** — `peak()` excludes  $\mathbf{q} = 0$  for that reason, not as a plotting convenience.

Validated against a closed form and against a quantity computed by an unrelated route. For free electrons  $N(q) = \Omega/4\pi^2 q$  below  $2k_F$  and zero above: the code reproduces the  $1/q$  law to  $10^{-4}$  and the hard cut to  $10^{-16}$ . A half-filled hydrogen chain has  $2k_F = \pi/c$ , so it must nest at  $\mathbf{q} = (0, 0, \frac{1}{2})$  — it does, at **99.8%** of the Cauchy-Schwarz bound, and `relax_spiral_q` relaxes the corresponding spin spiral to 0.500014 by going downhill in a magnetic total energy, which shares nothing with this. Against the vendored Elk binary on aluminium,  $N(0)/D(E_F)^2$  agrees to 1.5%; the raw  $N(0)$  differs by 4.7%, which is the all-electron Fermi surface underneath and is quoted rather than tuned away.

**Refuses.** A constrained `tot_magnetization`, which has one Fermi level per channel, so  $g(\mathbf{k})$  is two different surfaces; a **fixed-occupation** run, which never searches for a level, so  $\delta(\varepsilon - E_F)$  has no argument; and a **spin spiral**, because the quantity is a statement about the state a spiral grows out of — compute it on the paramagnet and compare its peak with `relax_spiral_q`'s answer. The delta defaults to a **Gaussian** even when the run used Methfessel-Paxton or cold smearing: those go negative on the wings, so  $g(\mathbf{k})$  can be negative and  $N(\mathbf{q})$  would have no sign at all. Pass `smearing=` to have the run's own anyway.

## 10.4 Magnetocrystalline anisotropy

The energy it costs to point a magnet's moment one way rather than another — what makes a hard magnet hard and what pins a spiral into a plane. It is tenths of a meV per atom against total energies of hundreds of Rydberg, so taking it as a difference of two self-consistent total energies asks two SCF runs to agree in their ninth digit. The *force theorem* does not ask that. Converge the magnet **without** spin-orbit coupling, rotate the converged density so its magnetization points along  $\mathbf{n}$ , and diagonalise **once** with the coupling switched on. At frozen density every term of the total energy except the band energy is a functional of  $\rho$  alone, so

$$E(\mathbf{n}_1) - E(\mathbf{n}_2) = E_{\text{band}}(\mathbf{n}_1) - E_{\text{band}}(\mathbf{n}_2), \quad E_{\text{band}} = \sum_{n\mathbf{k}} w_{n\mathbf{k}} \varepsilon_{n\mathbf{k}},$$

exactly, and the whole calculation is one diagonalisation per direction.

**The cheaper thing gives zero.** Freezing the wavefunctions as well and taking  $\langle \psi | H_{\text{SOC}} | \psi \rangle$  once returns no anisotropy at all: the coupling enters at first order as  $\xi \langle \mathbf{L} \rangle \cdot \mathbf{n}$ , and the orbital moment of a scalar-relativistic collinear state is quenched. The anisotropy is second order in the coupling, and what supplies it is the repulsion between levels that a diagonalisation performs and an expectation value does not. `frozen_expectation` computes that vanishing term anyway.

**Two pseudopotential files, one density.** The first leg runs with a scalar-relativistic dataset and the one-shot leg with the fully-relativistic dataset of the *same generation*;



only the density crosses, and a density does not know which file made it. That is what makes an ultrasoft anisotropy reachable, the alternative — averaging one relativistic file's  $j$  channels back — being refused for ultrasoft and PAW by `pw.x` itself.

```
from defumat import Calculator

scalar = Calculator.from_file("scf.in")          # nspin = 2, Co.pbe-nd-rrkjus
spinor = Calculator.from_file("nscf.in")        # noncolin, lspinorb, rel- file
mae = scalar.get_anisotropy(spinor, directions="xyz")

print(mae.anisotropy_mev, mae.easy_axis)
print(mae.difference(0, 2))                    # in-plane minus out-of-plane
```

`directions` takes a list of vectors, an integer asking for that many points spread evenly over the sphere, or one of "cardinal", "xz" and "xyz". `projected=True` adds the decomposition of the band energy over atomic orbitals, resolved by atom,  $\ell$ ,  $m$  and spin, which is what says *which* orbitals supply the anisotropy; comparing its total against  $E_{\text{band}}$  measures the spilling.

`soc_scale` scales the spin-orbit part of the nonlocal potential and the overlap while keeping the same dataset and the same k-points. At 0 the anisotropy is exactly zero, which is the control for the whole calculation: without the coupling the Hamiltonian is invariant under a global spin rotation, so a band energy *cannot* depend on where the moment points.

**The torque: the same number from one angle.** Taking the anisotropy as  $E(\mathbf{n}_1) - E(\mathbf{n}_2)$  is  $10^{-5}$  Ry out of  $10^2$  — seven digits of cancellation. The torque takes it as a *derivative* evaluated once, where nothing cancels: for  $E(\theta) = K_1 \sin^2 \theta$ ,

$$-\frac{dF}{d\theta} = -K_1 \sin 2\theta, \quad \text{so at } \theta = 45^\circ \text{ the torque is } -K_1.$$

```
torque = scalar.get_torque(spinor)              # 45 degrees by default
print(torque.anisotropy_constant_mev)          # K1 in meV
```

**It differentiates the free energy, and for a smeared metal that is not the band energy.** A Hellmann-Feynman derivative at frozen occupations gives  $dF/d\theta$  with  $F = \sum w\varepsilon - TS$ ;  $\sum w\varepsilon$  carries an extra  $\sum (dw/d\theta)\varepsilon$  that the entropy cancels. The difference is not small. On tetragonal cobalt, as the smearing narrows:

| degauss (Ry) | band difference | free difference | torque |
|--------------|-----------------|-----------------|--------|
| 0.020        | 1.2353          | 0.5523          | 0.5523 |
| 0.005        | 0.5434          | 0.5291          | 0.5291 |
| 0.002        | 0.5308          | 0.5308          | 0.5308 |

all in meV. The torque tracks the free-energy difference at every width and both converge on  $K_1 = 0.531$ ; at a production degauss of 0.02 the *band*-energy difference reads 2.3 times that. Use `MagneticAnisotropy.free_energies` to compare like with like.

**Refuses. PAW**, because the handoff from the collinear run carries the density and nothing else, and a PAW Hamiltonian needs `ddd_paw`, which is built from a `becsum` belonging to a run with a different pseudopotential file (`pw.x` refuses it in the same place). A **Hubbard**  $U$ , whose `ns` is not in the handoff either; a potential-only **meta-GGA**,



whose  $\tau$  is not; a converged **magnetic field** or constrained moment, whose energy sits outside the reported total; and a **spin spiral**, which refuses spin-orbit coupling permanently and so has no coupling to switch on. A direction other than the system's own `angle1/angle2` needs `nosym`: a magnetic noncollinear run reduces its k-grid with the *magnetic* symmetry group, which depends on where the moment points, so two directions would otherwise be sampled on two different wedges. `soc_scale` takes 0 or 1 and nothing in between: the coupling-free end of the overlap is spin-independent, so the anisotropy vanishes identically there whatever else it is, but a blend of it with the true overlap partway is not the overlap of any set of projectors and  $S$  stops being a usable metric.

## 10.5 X-ray and magnetic structure factors

What a diffractometer measures is not a density but its Fourier coefficients, one per reflection:

$$F(\mathbf{H}) = \int_{\Omega} \rho(\mathbf{r}) e^{i\mathbf{H}\cdot\mathbf{r}} d^3r, \quad F_j(\mathbf{H}) = \int_{\Omega} m_j(\mathbf{r}) e^{i\mathbf{H}\cdot\mathbf{r}} d^3r$$

for every reciprocal lattice vector with  $|\mathbf{H}| \leq h_{\max}$  — X-rays scattering off the charge, neutrons off the magnetization. Elk's tasks 195 and 196; `pw.x` has neither. In a plane-wave code the integral is an array that already exists, up to the crystallographers' two conventions (the positive phase, and no  $1/\Omega$ ), so the whole step is one transform and a gather: **6 ms** against the seconds its SCF cost.  $F(0)$  is the electron count of the cell and  $F_j(0)$  its magnetic moment, which is what anchors both.

```
from defumat import Calculator

calc = Calculator.from_file("si.scf.in")
factors = calc.get_structure_factors(hmax=5.0)

factors.charge           # (nh,) complex F(H), in electrons
factors.magnetization     # (nh, ncomp) in Bohr magnetons, or None
factors.vectors.indices   # (nh, 3) the (h k l) after the output transform
factors.vectors.length    # (nh,) |H| in 1/bohr
factors.vectors.multiplicity # the size of each star that was collapsed
factors.of((2, -2, -2))   # one reflection, by the index the table prints
print(factors.table(limit=10))

# Elk's wsfac: rebuild the density from the states in an energy window (Ry),
# which is what makes this a probe of bonding rather than of the whole charge.
bonding = calc.get_structure_factors(hmax=3.0, window=(-1.0, 0.2))
```

The functional entry point is `defumat.workflows.sfac.run_structure_factors`. `transform="conventional"` labels the reflections on the conventional cubic axes (Elk's `vhmat`), `reduce=False` keeps every member of every star, and `core=True` adds a dataset's frozen core-correction charge.

**A pseudopotential density is valence-only, and the comparison with an all-electron code is structured by reflection class rather than by one tolerance.** On silicon, against the vendored Elk binary on the same cell and k-grid: the origin is 8 electrons here and 28 there, and every *allowed* reflection is core-dominated in the same proportion — (111) is 1.7495 against 15.14, and no tolerance makes that agree. What *is* comparable is a **forbidden** reflection, where the spherical part of every atom cancels and the core with it: silicon's (222) is **0.347406** here against **0.33416**, 4%, on a quantity whose entire value is the aspherical bonding charge. That last statement is checked without any other

code, since this package’s own SCF starting guess is a superposition of free atoms: it gives 0 at (222) to  $10^{-10}$ . The space-group forbidden reflections — diamond’s (200), (420) — vanish to  $10^{-9}$  here and  $10^{-13}$  there.

**Norm-conserving, ultrasoft and PAW all work, and the difference between them is measurable.**  $F(0) = 8.0000000000$  on each, which is the check that the augmentation charge reached the transform at all — for a soft dataset most of the density near the nucleus is that charge. The like-for-like reference is not Elk’s all-electron total but its **valence-only** run, which `wsfac` above the 2p core produces: against it the augmentation charge is worth a factor of two, (111) coming out  $-0.7\%$  for ultrasoft and PAW against  $-1.5\%$  norm-conserving, and (004)  $+31\%$  against  $+61\%$ . It is worth *nothing* on the forbidden (222), which is 0.3474 on all three to  $3 \times 10^{-4}$ : a silicon atom in diamond sits at  $43m$ , whose lowest non-spherical invariant is  $l = 3$ , and an  $s, p$  dataset’s augmentation charge carries multipoles only to  $L = 2$ . That is also what the 4% against all-electron is, by elimination — not the k-grid, not either basis, not the functional (0.07%), and **not the core**, since Elk’s own valence-only run gives the identical 0.334165.

**Magnetic reflections are a low- $|\mathbf{H}|$  quantity and the boundary is measured.** Ferromagnetic bcc iron against Elk’s task 196, comparing the normalized  $F_{\text{mag}}(\mathbf{H})/F_{\text{mag}}(0)$  because the moments differ ( $2.2145 \mu_B$  here,  $2.0613$  there): 5.8% at the first reflection, 17% at the second, worse outwards. Raising `ecutwfc` from 30 to 45 moves those by less than 0.1%, so it is the pseudisation and not the basis — iron’s moment is in the 3d shell, inside the radius the dataset smooths. The cheapest cell carries the sharpest magnetic statement instead: an antiferromagnetic hydrogen chain has charge factors 2.000000, 0, 0.8587, 0 along the chain and magnetic ones 0, 1.2254, 0, 0.3862, so it scatters neutrons exactly where it scatters no X-rays, and the zeros are exact.

**Refuses.** An `hmax` past  $\sqrt{\text{ecutrho}}$ : the FFT box carries a coefficient at every frequency out to its corner, and past the density’s own cutoff those are the aliased tail of  $|\psi|^2$  rather than Fourier components — small, smooth and entirely plausible, which is why this is a guard and not a check on the answer. An empty or inverted energy window, and `core=True` on a dataset with no nonlinear core correction. **Said rather than refused**, because the numbers are right and only a claim about them would be wrong: these are *valence* structure factors, and a star’s multiplicity applies to the charge, the magnetic members of a star being related by  $\det(R) R$  rather than being equal.

## 11 Things with no `pw.x` counterpart

**Spin spirals** by the generalized Bloch theorem: a spiral of wavevector  $\mathbf{q}$  is a spinor whose two components live on *different* plane-wave spheres,

$$\psi_{\mathbf{k}}(\mathbf{r}) = \begin{pmatrix} e^{i(\mathbf{k}-\mathbf{q}/2)\cdot\mathbf{r}} u^{\uparrow}(\mathbf{r}) \\ e^{i(\mathbf{k}+\mathbf{q}/2)\cdot\mathbf{r}} u^{\downarrow}(\mathbf{r}) \end{pmatrix}$$

so in the rotated frame the density is lattice periodic and the SCF, the functional and the mixer are untouched. Only two terms carry  $\mathbf{q}$  —  $|\mathbf{k} \mp \mathbf{q}/2 + \mathbf{G}|^2$  and  $v_{\text{nl}}$  — and  $dE/d\mathbf{q}$  at frozen coefficients is `jax.grad` of them.

```
from defumat.workflows.spiral import (run_spiral_scan, relax_spiral_q,
                                     heisenberg_exchange)
```

```

scan = run_spiral_scan(system, pseudos, wavevectors) # E(q)
best = relax_spiral_q(system, pseudos) # BFGS on the reciprocal metric
print(best.wavevector, best.scf.total_energy)

J = heisenberg_exchange(scan, system.cell, shells) # fit E(q) for the J(R)

scan = run_spiral_scan(system, pseudos, wavevectors, gradients=True)
print(scan.relative) # E(q) - E(q_0) in mRy, from the energies
print(scan.integrated) # the same curve, from dE/dq along the path

```

heisenberg\_exchange fits  $E(\mathbf{q}) - E(0) = m^2 \sum_{\mathbf{R}} J(\mathbf{R}) [1 - \cos(\mathbf{q} \cdot \mathbf{R})]$  over a set of lattice shells, which is how a spiral scan becomes a spin model. compute\_spiral\_gradient gives  $dE/d\mathbf{q}$  on its own.

**$E(\mathbf{q})$  two ways.** With gradients=True the scan also takes  $dE/d\mathbf{q}$  at every point, and scan.integrated accumulates it along the path,

$$E(\mathbf{q}_n) - E(\mathbf{q}_0) = \sum_m \int d\mathbf{q} \cdot \frac{\partial E}{\partial \mathbf{q}},$$

by the trapezoid rule on the segments between successive wavevectors (both factors in lattice coordinates, so the contraction needs no metric and the path may bend). The two curves are the same physics and the gap between them is the useful part: the energies are computed on a plane-wave sphere *rebuilt* at every point and step by the Pulay error of a finite basis wherever a plane wave crosses the cutoff, while the gradient is taken at a *frozen* sphere and does not see those steps. On the hydrogen chain of tests/data/qe/h-chain-spiral.in at ecutwfc = 25 the direct energies rise, fall and rise again over a 13-point scan where the integrated curve descends monotonically; the two agree to 0.005 mRy at ecutwfc = 60 once the quadrature error is removed, against 0.130 at 25.

Two error sources live in that gap and refining tells them apart: the trapezoid rule's own  $h^2$  shrinks when the  $\mathbf{q}$  path is refined (0.051  $\rightarrow$  0.016  $\rightarrow$  0.003 mRy for 7  $\rightarrow$  13  $\rightarrow$  25 points) and the basis noise does not (0.139, 0.138, 0.137). What the integrated route does *not* buy is worth stating, since the intuition that it might is a common one: it is not cheaper (every point still needs its own SCF), it does not converge on a coarser  $k$ -mesh (the gradient is the exact derivative of the *same* fixed-mesh energy, so it inherits the same Brillouin-zone error), and it needs a *tighter* conv\_thr rather than a looser one, an energy's error being second order in the density's where a derivative's is first.

Also here: **per-atom magnetic fields**, the **topological invariants** of the previous section, and **rank- $n$  tensor symmetrisation** — symme.f90's symmatrix3/symtensor3 written at any rank.

**Refuses.** A spiral with spin-orbit coupling, permanently — it breaks the theorem, and Elk refuses it too; a spiral with symmetry, until the spin space group is written, so a spiral needs nosym; a spiral with ultrasoft/PAW; and spiral relaxation under a magnetic field, whose energy is outside the reported total, so the state is stationary for a different functional than the one being differentiated.

## 12 Performance: dials, memory and GPUs

**Two batching dials** control how much is in flight, and both defaults follow the platform: on a CPU, QE's loop — one  $k$ -point and one band at a time, which is what a *cache* wants —

and on a GPU the whole of both axes, which is what a card wants. Either end is reachable everywhere:

```
DEFUMAT_K_BATCH=all DEFUMAT_BAND_BATCH=all python3 run.py
```

or per call, `run_scf(..., k_batch=...)`. The chunk size changes the order contributions are summed in and nothing else — round-off,  $\sim 1e-15$  Ry.

**Why the default is per platform.** A GPU has no cache to fit in and inverts the conclusion: on a ten-atom metal,  $k=1, b=1$  costs **4.5x** what  $k=all, b=all$  does on the same card, sixteen-atom silicon runs at **0.20x** — an outright loss against four CPU cores — and at thirty-two atoms  $b=1$  did not finish at all. So a run on a device gets both batched unless it says otherwise, and the plausible half-measure  $k=all, b=1$  is never the default: it is the *worst* setting available, because it buys the batched mode's memory with the looped mode's kernel launches. An explicit argument beats the environment variables, which beat the platform.

**Measured GPU speedups** (one H200 against one CPU core, per SCF iteration): 2x on a cheap metal, 3-5x on norm-conserving silicon, 9-21x with an augmentation charge or a magnetization, and **83-115x on spin-orbit**. Energies agree to  $\sim 1e-9$  Ry — the same agreements the CPU already has.

**Memory.** Spin-orbit is where it bites: a 20-atom bismuth cell with a relativistic ultrasoft dataset peaks at **34.7 GB**, which fits a 141 GB card and does not fit a 32 GB one. Everything else in that sweep sits under 2 GB.

**A response costs memory according to how it is differentiated**, which is worth knowing before starting one. A *forward* response — the Sternheimer solves behind  $\epsilon$  and  $Z^*$  — costs 0.7-1.0 GB over the SCF that produced the state, and a *forward-over-reverse* one — the dynamical matrix, the Raman tensors — 1-2 GB. The *stress* is the exception and it is a *reverse* derivative through the radial transforms: **10.5 GB** on eight-atom ultrasoft silicon, ten times any of the others, and an order more than the SCF it came from. Measure your own with `tools/gpu/phase5.py <case> --property <name>`.

**Precision.** Every dtype comes from `config.Precision`; x64 is enabled before any array exists, and all validation is float64. float32 is a performance mode and **never one a correctness claim is made in**.

## 13 Accuracy summary

Against Quantum ESPRESSO on the same input:

| quantity                                | agreement                          |
|-----------------------------------------|------------------------------------|
| total energy, norm-conserving LDA       | $\sim 1\text{e-9 Ry}$              |
| ultrasoft, PAW                          | $\leq 3\text{e-9 Ry}$              |
| PBE / revPBE / PBEsol                   | $\leq 6\text{e-9 Ry}$              |
| metals, every smearing and tetrahedra   | $2.5\text{e-8 Ry}$                 |
| collinear spin                          | $1.2\text{e-9 Ry}$                 |
| noncollinear magnetism                  | $2.8\text{e-9 Ry}$                 |
| spin-orbit coupling                     | $\leq 1.3\text{e-8 Ry}$            |
| DFT+U                                   | $\leq 6.7\text{e-9 Ry}$            |
| band structure                          | $0.0002\text{ eV}$                 |
| forces                                  | $\leq 2\text{e-5 Ry/bohr}$         |
| stress (13 cases)                       | $\leq 2.7\text{e-7 Ry/bohr}^3$     |
| relaxed geometry                        | $1\text{e-6 bohr}$                 |
| dielectric constant (NC, US, PAW)       | $\leq 1.2\text{e-4}$               |
| Born effective charges (NC)             | every printed digit                |
| phonons at $\Gamma$ , silicon           | $0.05\text{ cm}^{-1}$              |
| phonons at $\Gamma$ , silicon (US, PAW) | $0.02\text{--}0.03\text{ cm}^{-1}$ |
| phonons at $\Gamma$ , aluminium (metal) | $0.002\text{ cm}^{-1}$             |
| Raman/IR spectra vs dynmat.x            | every printed digit                |

**One case does not close, and is recorded rather than explained.** `bi10-soc` — ten bismuth atoms with a fully-relativistic dataset — sits **1.9e-4 Ry** from QE on a total of  $-1477.737$ , which is  $1.3\text{e-7}$  relative. Both codes choose the same symmetry operations, the same  $k$ -point, the same grid and the same  $G$ -vectors, and both converge. It is in the test suite at that tolerance, not hidden.

## Where to go next

- `README.md` — the short tour and a first calculation
- `notebooks/` — 26 worked examples on concrete systems, executed and committed, each with a `.md` export beside it
- `PERFORMANCE.md` — the running performance log, CPU and GPU
- `PLAN.md` — the architecture, the phase breakdown, the trap catalogue
- `GPU.md` — the GPU roadmap